

## Contents

Labs.....	3
Lab 1 – One param models, simulation, direct numerical evaluation.....	3
Assignment 1 – Bernoulli posterior and log odds posterior with convergence .....	3
Assignment 2 – Lognormal distribution, posterior vs real distrib, ginicoefficient, cred/HPD intervals.....	4
Assignment 3 – von Mises distribution with Bessel, plot posterior (normalization), finding posterior mode.....	6
Lab 2 – Polynomial regression, and classification with logistic regression.....	8
Assignment 1 – Setting proper priors, simulating from joint posterior, marginal posterior, scatter plot overlaid with posterior median and credible intervals, posterior mode prediction ..	8
Assignment 2 – Normal approx. through optimization, comparing with maximum likelihood, simulating from predictive distribution, multiple trials (binomial) .....	12
Lab 3 – MCMC using Gibbs sampling and Metropolis Hastings .....	15
Assignment 3 – Gibbs sampler with normal data, evaluate convergence, mixture of normals, plot of histogram kernel density estimate & normal density & mixture of normal .....	15
Assignment 2 – Linear regression with Poisson, ML-estimator, optim with Poisson, simulating through RWM Hastings, myFunction(...), simulate from predictive distribution .....	20
Lab 4 – HMC with Stan .....	24
Assignment 1 – AR(1)-process, stanmodel with AR(1)-process, Poisson conditioned on AR(1)-process in stan, plot with data & posterior mean & credible intervals over time .....	24
Exams.....	30
2017-05-30.....	30
Assignment 1 – Plot posterior (rice function), normal approx. through optim, simulation for new obs .....	30
Assignment 2 – Model posterior data, plot posterior and compare with data, Gibbs for mixture of models, graphical methods for evaluating.....	31
Assignment 3 – Linear regression (cars), marginal distributions, interpretation of cred. Interval, predictive simulation .....	34
Assignment 4 – Maximizing posterior expected utility.....	37
2017-08-16.....	38
Assignment 1 – Plot posterior density (Cauchy), normal approx. through optimization, marginal posterior.....	38
Assignment 2 – Bayes Linear Regression, decision with loss function, predictive distribution ...	39
Assignment 4 – Simulation from posterior (Poisson with Gamma prior), simulation predictive distribution, maximizing posterior utility.....	41
2017-10-27 .....	43

Assignment 1 – Plot posterior distrib (beta symmetric prior, exponential likelihood), joint posterior distribution through optim, how to choose model? .....	43
Assignment 2 – Bayes Linear Regression, HPD interval, predictive distribution house prices for specific house type .....	44
Assignment 4 – Simulation of predDraw from normal, approximated probability weight larger than value, decision making (linear loss function) .....	46
2018-06-01 .....	47
Assignment 1 – Plot posterior using samples and expression, simulate from predictive .....	47
Assignment 2 – Bayes Linear Regression (fish), marginal posterior, 90 % equal tail with interpretation, new experiment with two different new_obs with likelihoods (Bayesian analysis) .....	48
Assignment 3 – Choosing between 3 models, marginal likelihood .....	51
Assignment 4 – Truncated normal distrib, stan with time series model, plot of data & posterior mean & 95 % credible intervals over time .....	51
2019-08-21 .....	54
Assignment 1 – Bayes Linear Regression, point estimates, 95 % tail intervals, posterior mode and HPD 90 % intervals, predictive Bayesian analysis .....	54
Assignment 2 – Predictive draw earthquakes (math calculations as basis).....	57
Assignment 3 – Calc unnormalized posterior and plot normalized posterior .....	57
Assignment 4 – Simulate using RWM Hastings, simulation using metropolis hasting with gamma proposal density function, suggestions how to improve sampler, traceplots.....	58
2019-10-31 .....	60
Assignment 1 – Expected utility, calculations from math.....	60
Assignment 2 – Poisson likelihood with gamma prior, plot posterior, separation of data, now two indep poisson models, comparison .....	61
Assignment 3 – Simulation of joins posteriors with Gibbs sampling, traceplots .....	62
Assignment 4 – Stan, plot scatter plot & mean of posterior predictive distrib & 90 % equal tail intervals, stanmodel with heteroscedastic variance (different variance over time).....	62

## Labs

### Lab 1 – One param models, simulation, direct numerical evaluation

#### Assignment 1 – Bernoulli posterior and log odds posterior with convergence

*## Assignment 1: Let  $y_1, \dots, y_n$  be Bernoulli distributed with parameter  $\theta$ . Assume that you have obtained a sample  $y$  with  $s=5$  successes in  $n=20$  trials. Assume a  $\text{Beta}(\alpha_0, \beta_0)$  prior for  $\theta$  and let  $\alpha_0=\beta_0=2$ .*

*## a) Draw random numbers from the posterior  $\theta$  given  $y \sim \text{Beta}(\alpha_0+s, \beta_0+f)$  and verify graphically that the posterior mean and standard deviation converges to the true values as the number of random draws grows large.*

```
set.seed(12345)
alpha0=2
beta0=2
s=5
f=15
n=20

# Function for calculating the mean of a beta-distribution
calcMeanBeta = function(alpha, beta) {
  return(alpha/(alpha+beta))
}

# Function for calculating the standard deviation of a beta-distribution
calcStdDevBeta = function(alpha, beta) {
  return(sqrt(alpha*beta/((alpha+beta)^2*(alpha+beta+1))))
}

# Function for calculating the mean squared error of drawn data
calcMSE = function(n, mean, data){
  return(sqrt(1/(n-1)*sum((data-mean)^2)))
}

# Function for drawing random values from the betadistribution
drawBetaValues = function(n, alpha, beta) {
  return(rbeta(n, alpha, beta))
}

MeanOfPosterior = calcMeanBeta(alpha0+s, beta0+f)
StdOfPosterior = calcStdDevBeta(alpha0+s, beta0+f)
nVector = seq(1, 5000, 1)
meanVector=c()
stdVector=c()
for (i in nVector) {
  set.seed(12345)
  betaValues= drawBetaValues(i, alpha0+s, beta0+f)
  meanVector=c(meanVector, mean(betaValues))
  stdVector=c(stdVector, calcMSE(i, mean(betaValues), betaValues))
}
plot(nVector, meanVector, main="Plot of how the mean converges with respect to the number of draws")
```

```

t to number of draws",
  xlab="Number of draws", ylab="Mean", type="l")
abline(h=MeanOfPosterior, col="red")
plot(nVector, stdVector, main="Plot of how the standard deviation converge
s with respect to the number of draws",
  xlab="Number of draws", ylab="Standard deviation", type="l")
abline(h=StdOfPosterior, col="red")
## As seen in the plot the posterior mean as well as the posterior standar
d deviation converges towards its true
## value of approx 0.29 and 0.09 respectively as the number of randow draw
s grows large.

## b) Use simulation (nDraws=10000) to compute the posterior probability P
r(theta>0.3 given y) and compare with
## with the exact value

trueProb=1-pbeta(0.3, alpha0+s, beta0+f)
set.seed(12345)
draw10000=rbeta(10000, alpha0+s, beta0+f)
probHat=sum(draw10000>0.3)/10000
print(trueProb)
print(probHat)

## As seen in the results from both calculations the probHat is very close
to the true probability from the beta
## distribution. As the number of draws increases the approximated probabi
lity will converge towards the true
## value.

## c) Compute the posterior distribution of the log-odds phi= Log(theta/(1
-theta)) by simulation (nDraws=10000)

phi=log(draw10000/(1-draw10000))
hist(phi, breaks=20, main="Distribution of the log-odds")
plot(density(phi), main="Density function of phi")

```

Assignment 2 – Lognormal distribution, posterior vs real distrib, ginicoefficient, cred/HPD intervals

```

## Assignment 2: Assume that you have asked 10 randomly selected persons a
bout their monthly
## income(inthousandsSwedishKrona)andobtainedthefollowingtenobservations:
44, 25, 45, 52, 30, 63, 19, 50, 34
## and 67. A common model for non-negative continuous variables is the log
-normal distribution. The log-normal
## distribution Log(N(my, sigma^2)) has density function ... for y > 0, my
> 0 and sigma > 0. The log-normal
## distribution is related to the normal distribution as follows: if y ~ L
og N(my, sigma^2) then
## log y ~ N(my, sigma^2). Let y1,...,yn given my and simga^2 ~ Log N(my,
sigma^2), where my=3.7 is assumed to be
## known but sigma^2 is unknown with noninformative prior p(sigma^2) is pr
oportional to 1/sigma^2. The posterior
## for sigma^2 is the Inv - chitwo distribution with X(n, thao^2) distribu

```

tion, where  $\text{thao}^2 = \sum((\log(y_i) - \text{my})^2)/n$

*## a) Simulate 10 000 draws from the posterior of  $\sigma^2$  (assuming  $\text{my}=3.7$ ) and compare with the theoretical  
## with the theoretical Inv - chitwo distribution with  $X(n, \text{thao}^2)$  posterior distribution.*

```
library(geoR)
x=c(44, 25, 45, 52, 30, 63, 19, 50, 34, 67)
n=length(x)
my=3.7

#Function for calculating thao^2
calcThao = function(data, my, n) {
  return(sum((log(data)-my)^2)/n)
}

invchisquare <- function(x, df, taosq){
  first = ((taosq*df/2)^(df/2))/gamma(df/2)
  second = (exp((-df*taosq)/(2*x)))/(x^(1+df/2))
  return(first*second)
}

thaosq=calcThao(x, my, n)
set.seed(12345)
drawX=rchisq(10000, n)
sigmasq=(n)*thaosq/drawX
xvals=seq(0.001, 3, 0.001)
plot(density(sigmasq), main="Density of simulated sigma^2, black = simulated distrib., red = actual distrib.")
lines(xvals,invchisquare(xvals, n, thaosq), col="red")
```

*## As seen in the plot the theoretical distribution (red line) follows the simulated one with good precision. This  
## indicates that the simulation has been made correctly.*

*## b) The most common measure of income inequality is the Gini coefficient,  $G$ , where  $0 \leq G \leq 1$ .  $G=0$  means a  
## completely equal income distribution, whereas  $G=1$  means complete income inequality. See Wikipedia for more  
## information. It can be shown that  $G=2 \cdot \text{CDF-normal}(\sigma/\sqrt{2}) - 1$  when income follow a Log  $N(\text{my}, \sigma^2)$   
## distribution. Use the posterior draws in a) to compute the posterior distribution of the Gini coefficient  $G$   
## for the current data set.*

```
G=2*pnorm(sqrt(sigmasq/2), mean=0, sd=1)-1
hist(G, breaks=100)
plot(density(G), main="Density function of simulated values of the Gini coefficient")
```

*## As seen in the plot the gini coefficient is centered at around 0.2 which means a rather unequal distribution.*

```
## c) Use the posterior draws from b) to compute a 90% equal tail credible
interval for G. A 90% equal tail interval
## (a,b) cuts off 5% percent of the posterior probability mass to the left
of a, and 5% to the right of b. Also,
## do a kernel density estimate of the posterior of G using the density fu
nction in R with default settings,
## and use that kernel density estimate to compute a 90% Highest Posterior
Density interval for G. Compare the
## two intervals.
```

```
GSorted=sort(G)[(0.05*length(G)+1):(0.95*length(G))]
# 90 % credible interval for G through the simulated draws
G_CredInterval=c(min(GSorted),max(GSorted))
print(G_CredInterval)
plot(density(G), main="Density function of simulated values of the Gini co
efficient with credible intervals")
abline(v = G_CredInterval[1], col="blue")
abline(v = G_CredInterval[2], col="blue")
```

```
GDensity=density(G)
GDensity.df=data.frame(x=GDensity$x, y=GDensity$y)
GDensity.df=GDensity.df[order(-GDensity.df[,2]),]
index=dim(GDensity.df)[1]
GDensity.df$y=cumsum(GDensity.df$y)/sum(GDensity.df$y)
GDensity_CredInterval_Vals=GDensity.df[GDensity.df$y<0.90,]
GDensity_CredInterval=c(min(GDensity_CredInterval_Vals$x), max(GDensity_Cr
edInterval_Vals$x))
print(GDensity_CredInterval)
abline(v = GDensity_CredInterval[1], col="red")
abline(v = GDensity_CredInterval[2], col="red")
title(sub="Blue = Simulated credible interval, Red = Kernel estimated cred
ible interval")
```

## As seen in the plot the credible intervals are quite similar with small deviations.

Assignment 3 – von Mises distribution with Bessel, plot posterior (normalization), finding posterior mode

```
## Assignment 3: Bayesian inference for the concentration parameter in the
von Mises distribution. This exercise is concerned
## with directional data. The point is to show you that the posterior dist
ribution for somewhat weird models can be
## obtained by plotting it over a grid of values. The data points are obse
rved wind directions at a given location on
## ten different days. The data are recorded in degrees: (40, 303, 326, 28
5, 296, 314, 20, 308, 299, 296) where North
## is located at zero degrees (see Figure 1 on the next page, where the an
gles are measured clockwise). To fit with
## Wikipedias description of probability distributions for circular data w
e convert the data into radians -pi<=y<=pi.
## The 10 observations in radians are (-2.44,2.14,2.54,1.83,2.02,2.33,-2.7
9,2.23,2.07,2.02).
## Assume that these data points are independent observations following th
```

```

e von Mises distribution
##  $p(y \text{ given } \mu, k) = \exp(k \cdot \cos(y - \mu)) / (2 \cdot \pi \cdot I_0(k))$ ,  $-\pi \leq y \leq \pi$ , where  $I_0(k)$ 
is the modified Bessel function of the
## first kind of order zero (see ?besselI in R). The parameter  $\mu$  ( $-\pi \leq \mu \leq \pi$ ) is the mean direction and  $k > 0$  is
## called the concentration parameter. Large  $k$  gives a small variance around  $\mu$ , and vice versa. Assume that  $\mu$  is
## known to be 2.39. Let  $K \sim \text{Exponential}(\text{Lambda}=1)$  a priori, where  $\text{Lambda}$ 
is the rate parameter of the exponential
## distribution (so that the mean is  $1/\text{Lambda}$ ).

## a) Plot the posterior distribution of  $k$  for the wind direction data over a fine grid of  $k$  values.

data_radian=c(-2.44,2.14,2.54,1.83,2.02,2.33,-2.79,2.23,2.07,2.02)
my=2.39
lambda=1

# Function for computing the vonMisesDistrib for a given dataset
vonMisesDistrib = function(kappa, data, my){
  likelihood=1
  for (i in data) {
    likelihood=likelihood*exp(kappa*cos(i-my))/(2*pi*besselI(kappa, 0))
  }
  return(likelihood)
}

# Function for computing the exponential distribution
exponDistrib = function(data, lambda) {
  return(lambda*exp(-lambda*data))
}

kappa_values=seq(0,10,0.01)

# Function for computing the posterior distribution
posteriorDistrib = function(kappa, lambda, data, my) {
  likelihood=vonMisesDistrib(kappa, data, my)
  prior=exponDistrib(kappa, lambda)
  return(likelihood*prior)
}

posteriorLikelihood=posteriorDistrib(kappa_values, lambda, data_radian, my)
posterior.df=data.frame(kappa=kappa_values, likelihood=posteriorLikelihood)
sumOfPosterior=sum(posterior.df$likelihood)
posterior.df$likelihood=posterior.df$likelihood*(1/sumOfPosterior)
final_sum=sum(posterior.df$likelihood)
plot(kappa_values, posterior.df$likelihood, xlab="Kappa", ylab="Likelihood",
     main="Posterior likelihood for different kappavalues", type="l", col="blue")

## As seen in the plot the Likelihood of the posterior peaks between 2 and

```

4 and then dies off for larger  
## kappa-values.

## b) Find the (approximate) posterior mode of  $k$  from the information in a ).

# Puts likelihood values with corresponding kappa-values to be able to retrieve the kappa-value corresponding to  
## the highest likelihood (mode)

```
posteriorMode=subset(posterior.df, likelihood==max(likelihood), kappa)
print(posteriorMode$kappa)
```

## The approximated posterior mode is found to be 2.12.

## Lab 2 – Polynomial regression, and classification with logistic regression

Assignment 1 – Setting proper priors, simulating from joint posterior, marginal posterior, scatter plot overlaid with posterior median and credible intervals, posterior mode prediction

## Assignment 1: The dataset TempLinkoping.txt contains daily average temperatures (in Celcius degrees) at

## Malmslatt, Linkoping over the course of the year 2018. The response variable is temp and the covariate is

## time=(the number of days since beginning of year)/365

## Your task is to perform a Bayesian analysis of a quadratic regression

##  $\text{temp} = \beta_0 + \beta_1 \cdot \text{time} + \beta_2 \cdot \text{time}^2 + \epsilon$ ,  $\epsilon \sim N(0, \sigma^2)$

## a) Determining the prior distribution of the model parameters. Use the conjugate prior for the linear

## regression model. Your task is to set the prior hyperparameters  $\mu_0$ ,  $\omega_0$ ,  $\nu_0$  and  $\sigma_0^2$  to sensible

## values. Start with  $\mu_0 = (-10, 100, -100)^T$ ,  $\omega_0 = 0.01 \cdot I_3$ ,  $\nu_0 = 4$  and  $\sigma_0^2 = 1$ .  $\theta = 1$ . Check if this prior

## agrees with your prior opinions by simulating draws from the joint prior of all parameters and for every draw

## compute the regression curve. This gives a collection of regression curves, one for each draw from the prior.

## Do the collection of curves look reasonable? If not, change the prior hyperparameters until the collection

## of prior regression curves agrees with your prior beliefs about the regression curve. [Hint: the R package

## mvtnorm will be handy. And use your Inv-chisquared simulator from Lab1.

# Read file

```
temp = read.table("TempLinkoping.txt", header=TRUE)
```

```
## install.packages("mvtnorm")
```

```
library(mvtnorm)
```

# Defining the parameters for the prior distribution

# Switched to  $\beta_0 = 0$  since it seems more reasonable and -10 seems too low  
.

```
my0=c(-10,100,-100)
```

```
omega0=0.5*diag(3)
```

# Using  $\nu_0 = 365$  since we have 365 observations



```

v0=365
sigma0_sq=0.5
omega0Inv=solve(omega0)

# Function for returning the response variable
calcRegr = function(betaMatrix, row, x) {
  return(betaMatrix[row,1]+betaMatrix[row,2]*x+betaMatrix[row,3]*x^2)
}

# Function for drawing simulated betavalues
drawBeta = function(my, sigma_sq, omegaInv) {
  return(rmvnorm(1, mean=my, sigma=sigma_sq*omegaInv))
}

nDraws=1000
set.seed(12345)
drawX=rchisq(nDraws, v0)
sigma_sq=(v0)*sigma0_sq/drawX
betaMatrix=matrix(0,nDraws,3)
# Create new plot with specific settings so that the loop can overlay plot
s
plot.new()
plot.window(xlim=c(0,1), ylim=c(-50,50))
axis(side=1)
axis(side=2)
set.seed(12345)
for (i in 1:nDraws) {
  betaMatrix[i,]=drawBeta(my0, sigma_sq[i], omega0Inv)
  lines(temp$time, calcRegr(betaMatrix, i, temp$time), col=rgb(0,0,0,0.2))
}
title(main="Temps depending on different times for different simulated mod
els", xlab="Time", ylab="Temp")

## The collection of curves look reasonable and in line with our prior bel
iefs. The temperature rises during the
## summer months and stays low in the beginning and the end of the year re
spectively.However, the value of -10
## were switched to 0 since it seems more reasonable with a measurement of
the temperature 0 on the 1st of
## January than a measurement of -10.

## b) Write a program that simulates from the joint posterior distribution
of beta0, beta1, beta2 and sigma^2.
## Plot the marginal posteriors of each parameter as a histogram. Also pro
duce another figure with a scatter plot
## of the temperature data and overlay a curve for the posterior median of
the regression function
## f(time)=beta0+beta1*time+beta2*time^2, computed for every value of time
. Also overlay curves for the lower
## 2.5% and upper 97.5% posterior credible interval for f(time). That is,
compute the 95% equal tail posterior
## probability intervals for every value of time and then connect the Lowe
r and upper limits of the interval by
## curves. Does the interval bands contain most of the data points? Should

```

they?

```
# Calculating the parameters for the posterior distribution
v_n=v0+length(temp$temp)
X=cbind(1, temp$time, temp$time^2)
Y=temp$temp
beta_hat=solve(t(X)%*%X)%*%t(X)%*%Y
my_n=solve(t(X)%*%X+omega0)%*%(t(X)%*%X)%*%beta_hat+omega0)%*%my0)
omega_n=t(X)%*%X+omega0
omega_n_Inv=solve(omega_n)
sigma_sq_n=(v0*sigma0_sq+(t(Y)%*%Y+t(my0)%*%omega0)%*%my0-t(my_n)%*%omega_n
)%*%my_n))/v_n

# Simulate the joint posterior
sigma_sq_post=(v_n)*c(sigma_sq_n)/drawX
betaMatrix_post=matrix(0,nDraws,3)
response_post_temp=matrix(0,nDraws,length(temp$time))
for (i in 1:nDraws) {
  betaMatrix_post[i,]=drawBeta(my_n, sigma_sq_post[i], omega_n_Inv)
}
# Plots the marginal distributions for the different beta-values
hist(betaMatrix_post[,1], breaks=100, main="Marginal posterior for beta0")
hist(betaMatrix_post[,2], breaks=100, main="Marginal posterior for beta1")
hist(betaMatrix_post[,3], breaks=100, main="Marginal posterior for beta2")
hist(sigma_sq_post, breaks=100, main="Marginal posterior for sigmasq")

plot(temp$time, Y, main="Plot of the temp data for different times", col="
blue",
      xlab="Time coefficient", ylab="Temp")
# Applies function calcRegr to the time-values for each of the drawn betas
and stores the results in matrix
for (i in 1:nDraws) {
  betaTemp=sapply(temp$time, calcRegr, betaMatrix=betaMatrix_post, row=i)
  response_post_temp[i,]=betaTemp
}

response_post=c()
credInterval=matrix(0, length(temp$time), 2)
# Retrieves the median of the response values as well as obtaining the upp
er and lower bound of credInterval
for (i in 1:length(temp$time)) {
  sortedTemp=sort(response_post_temp[,i])
  response_post=c(response_post, (sortedTemp[500]+sortedTemp[501])/2)
  credInterval[i,]=quantile(response_post_temp[,i], probs=c(0.025, 0.975))
}

lines(temp$time, response_post)
lines(temp$time, credInterval[,1], lty=21, col="gray")
lines(temp$time, credInterval[,2], lty=21, col="gray")
title(sub="Grey = 95 % credible intervals, Black = Median")

## The interval bands contain most of the data points. They should contain
most of the data points if the model
## is accurate in terms of describing the reality. In this case, it seems
```

Like the model has captured most of  
## the data points which means that the model describes the reality fairly well.

## c) It is of interest to locate the time with the highest expected temperature (that is, the time where  
##  $f(\text{time})$  is maximal). Let's call this value  $x_{\text{tilde}}$ . Use the simulations in b) to simulate from posterior  
## distribution of  $x_{\text{tilde}}$ . [Hint: The regression curve is quadratic. You can find a simple formula for  $x_{\text{tilde}}$   
## given  $\beta_0$ ,  $\beta_1$  and  $\beta_2$ ]

```
# Function for calculating the time-value which yields the maximum response (the derivative of response function)
calcMaxTemp = function(betaMatrix, row) {
  return(-betaMatrix[row,2]/(2*betaMatrix[row,3]))
}
```

```
# For each of the draws the time-value which yields the maximum temperature is stored in a vector
time_max_temp=c()
for (i in 1:nDraws) {
  time_max_temp=c(time_max_temp, calcMaxTemp(betaMatrix_post, i))
}
```

```
hist(time_max_temp, breaks=10, xlim=c(0,1), main="Frequency of max temperatures simulated from  $x_{\text{tilde}}$ ",
      xlab="Temperature")
```

## As seen in the histogram the derived highest temperature from the simulated models is mostly present in late  
## June which seems reasonable if applying to Malmslätt in Sweden where the temperature is the highest during the  
## summer time.

## d) Say now that you want to estimate a polynomial model of order 7, but you suspect that higher order terms  
## may not be needed, and you worry about overfitting. Suggest a suitable prior that mitigates this potential  
## problem. You do not need to compute the posterior, just write down your prior. [Hint: the task is to specify  
##  $\mu_0$  and  $\omega_0$  in a smart way.]

## A suitable prior for this task would be to set  $\mu_0$  to 0 since you want most of the coefficients close to zero  
## to obtain increased shrinkage. You would also want to set  $\omega_0$  to  $\lambda \text{Id} * \text{IdentityMatrix}$ . This would mean  
## that for larger values of  $\lambda$  more and more of the beta values would be close to zero since the spread of  
## the distribution of the beta values would decrease. In this case, where there is a worry about overfitting,  
## it might be a good idea to choose a large  $\lambda$  to decrease the spread of the beta values and increase the  
## probability that most of the beta values are around 0.

Assignment 2 – Normal approx. through optimization, comparing with maximum likelihood, simulating from predictive distribution, multiple trials (binomial)

```
## Assignment 2: Consider the Logistic regression  $\Pr(y=1 \text{ given } x) = \exp(x^T \text{Beta}) / (1 + \exp(x^T \text{Beta}))$ , where  $y$  is the
## binary variable with  $y = 1$  if the woman works and  $y = 0$  if she does not
##  $x$  is a 8-dimensional vector containing
## the eight features (including a one for the constant term that models the
## intercept). The goal is to approximate
## the posterior distribution of the 8-dim parameter vector  $\beta$  with multivariate normal distribution.
##  $\text{Beta}$  given  $y$  and  $x \sim N(\tilde{\text{Beta}}, JY(\tilde{\text{Beta}})^{-1})$ , where  $\tilde{\text{Beta}}$  is the posterior mode and  $J(\tilde{\text{Beta}})$  is the second
## derivative, the observed Hessian evaluated at the posterior mode. It is
## actually not hard to compute this
## derivative by hand, but don't worry, we will let the computer do it numerically for you. Now, both  $\tilde{\text{Beta}}$  and
##  $J(\tilde{\text{Beta}})$  are computed by the optim function in R. I want you to implement an own version of my example code at
## the website. You can use my code as a template, but I want you to write your own file so that you understand
## every line of your code. Don't just copy my code. Use the prior  $\text{Beta} \sim N(0, \text{thao}^2 I)$  with  $\text{thao}=10$ . Your report
## should include your code as well as numerical values for  $\tilde{\text{Beta}}$  and  $JY(\tilde{\text{Beta}})^{-1}$  for the WomenWork data. Compute
## an approximate 95% credible interval for the variable NSmallChild. Would you say that this feature is an
## important determinant of the probability that a woman works? [Hint: To verify that your results are reasonable,
## you can compare to you get by estimating the parameters using maximum likelihood.
## glmmodel = glm(Work~0+., data=WomenWork, family=binomial)
```

```
# Use of libraries
```

```
library(mvtnorm)
```

```
# Read data
```

```
WomenWork = read.table("WomenWork.dat", header=TRUE)
```

```
# User input
```

```
tau = 10
```

```
# Defining vectors  $X$  and  $Y$ 
```

```
 $X = \text{as.matrix}(\text{WomenWork}[2:\text{ncol}(\text{WomenWork})])$ 
```

```
 $Y = \text{WomenWork}[1]$ 
```

```
 $n\text{Features} = \text{dim}(X)[2]$ 
```

```
 $\text{covNames} = \text{names}(\text{WomenWork}[2:\text{ncol}(\text{WomenWork})])$ 
```

```
# Constructing prior
```

```
 $\mu_{\text{prior}} = \text{rep}(0, n\text{Features})$ 
```

```
 $\sigma_{\text{prior}} = \text{tau}^2 * \text{diag}(n\text{Features})$ 
```

```
# Defining function for returning the log posterior
```

```

logPostLogistic = function(beta, Y, X, mu, sigma) {
  nFeat = length(beta)
  XBeta=X%%beta
  # Defining LogLikelihood
  logLike = sum(Y*XBeta-log(1+exp(XBeta)))
  # Defining prior
  prior = dmvnorm(beta, mean=mu, sigma=sigma, log=TRUE)
  # Adding Loglikelihood and Logprior together. Since it is log both of them
  # are added instead of multiplied
  return(logLike + prior)
}

# Defining initial values to be passed on to the optimizer
set.seed(12345)
initVals = rnorm(dim(X)[2])

# Finding the optimized betavector
optimResult = optim(initVals, logPostLogistic, Y=Y, X=X, mu=mu_prior, sigma=sigma_prior, method=c("BFGS"),
                    control=list(fnscale=-1), hessian=TRUE)

# Defining the values of interest
postMode = optimResult$par
postCov = -solve(optimResult$hessian)
names(postMode) = covNames
approx_PostStd = sqrt(diag(postCov))
names(approx_PostStd) = covNames
print("The posterior mode is:")
print(postMode)
print("The approximated standard deviations are:")
print(approx_PostStd)

# Compute marginal distribution for NSmallChild
NSmallChild_mode = as.numeric(postMode["NSmallChild"])
NSmallChild_std = as.numeric(approx_PostStd["NSmallChild"])
credInterval_NSsmallChild = qnorm(p=c(0.025, 0.975), mean=NSmallChild_mode,
sd=NSmallChild_std)
print(paste("The lower bound of the 95 % credible interval for the feature NSmallChild is",
            round(credInterval_NSsmallChild[1], 6), "and the upper bound is",
            round(credInterval_NSsmallChild[2], 6)))

# Control that the calculations have been made correctly
glmModel = glm(Work ~ 0+., data=WomenWork, family=binomial)
print(glmModel$coefficients)
print(postMode)

## Since the values for the credible interval for NSmallChild are quite large in the negative direction it is
## reasonable to conclude that the feature NSmallChild affects the response variable fairly much towards the
## response 0 which means that the woman doesn't work. This seems like a reasonable conclusion in terms of how

```

```
## it is in reality as well. When checking if the results are reasonable,
a comparison was made with an
## estimation using the maximum likelihood method. The results was very si
milar which strongly suggests that
## the results obtained from the code are reasonable.
```

```
## b) Write a function that simulates from the predictive distribution of
the response variable in a logistic
## regression. Use your normal approximation from 2(a). Use that function
to simulate and plot the predictive
## distribution for the Work variable for a 40 year old woman, with two ch
ildren (3 and 9 years old), 8 years
## of education, 10 years of experience. and a husband with an income of 1
0. [Hints: The R package mvtnorm will
## again be handy. Remember my discussion on how Bayesian prediction can b
e done by simulation.]
```

```
sigmoid = function(value) {
  return (exp(value)/(1+exp(value)))
}
```

```
makePredLogReg = function(data, mean, sigma, nDraws) {
  betaPred = rmvnorm(nDraws, mean=mean, sigma=sigma)
  linearPred = betaPred %*% data
  logPred = sigmoid(linearPred)
  return(logPred)
}
```

```
nDraws=10000
woman=c(1, 10, 8, 10, (10/10)^2, 40, 1, 1)
set.seed(12345)
womanWorkPred=makePredLogReg(woman, postMode, postCov, nDraws)
logistic_distrib=c()
for (i in womanWorkPred) {
  logistic_distrib=c(logistic_distrib, rbinom(1, 1, i))
}
barplot(table(logistic_distrib), main="Histogram of the predicted probabil
ities")
```

```
## As seen in the plots the calculated probabilities of the woman in quest
ion working is fairly low. The highest
## density is seen in the range between 0.2 and 0.3 approximately. This al
so makes sense if applied to a real
## situation. A woman with a small child is likely to stay at home with th
e child, i.e. not working. If the
## classification of the response variable results in "working" if the pre
dicted probabilities is above 0.5 and
## "not working" otherwise, it is clear from the distribution that the cla
ssification of a woman working, with
## the parameters inputted, is very unlikely to happen.
```

```
## c) Now, consider 10 women which all have the same features as the woman
in 2(b). Rewrite your function and
## plot the predictive distribution for the number of women, out of these
```



10, that are working.  
## [Hint: Which distribution can be described as a sum of Bernoulli random variables?]

```
makePredLogRegMultiple = function(data, mean, sigma, nDraws, n) {  
  multiplePred=c()  
  for (i in 1:nDraws) {  
    betaDraw = makePredLogReg(data, mean, sigma, 1)  
    multiplePred=c(multiplePred, rbinom(1, n, betaDraw))  
  }  
  barplot(table(multiplePred), main=paste("Distribution for prediction made on", n, "women"),  
          xlab="No. of women")  
}
```

```
makePredLogRegMultiple(woman, postMode, postCov, 10000, 10)
```

## As seen in the histogram the binomial case resembles the density of predicted probabilities with the  
## highest density found at 2 women. This result seems reasonable since when the number of draws taken from  
## the binomial distribution goes towards infinity the shape of the corresponding distribution will resemble  
## the shape of the distribution for the probability  $p$  in the Bernoulli case, more and more.

### Lab 3 – MCMC using Gibbs sampling and Metropolis Hastings

Assignment 3 – Gibbs sampler with normal data, evaluate convergence, mixture of normals, plot of histogram kernel density estimate & normal density & mixture of normal

## Assignment 1: The data rainfall.dat consist of daily records, from the beginning of 1948 to the end of 1983,  
## of precipitation (rain or snow in units of 1/100 inch, and records of zero precipitation are excluded) at  
## Snoqualmie Falls Washington. Analyze the data using the following two models.

## a) Assume the daily precipitation ( $y_1, \dots, y_n$ ) are iid normally distributed,  
##  $y_1, \dots, y_n$  given  $\mu$  and  $\sigma^2 \sim N(\mu, \sigma^2)$  where both  $\mu$  and  $\sigma^2$  are unknown. Let  $\mu \sim N(\mu_0, \tau_0^2)$   
## independently of  $\sigma^2 \sim \text{Inv chisquare}(\nu_0, \sigma_0^2)$   
## i) Implement (code!) a Gibbs sampler that simulates from the joint posterior  $p(\mu, \sigma^2 \text{ given } y_1, \dots, y_n)$ .  
## The full conditional posteriors are given on the slides from Lecture 7.

```
library(mvtnorm)  
# Read data  
Rainfall = read.table("rainfall.dat")  
  
# Setup  
# Prior knowledge of  $\mu_0$  taken from Google  
 $\mu_0=14.79$   
mean_rainfall=mean(Rainfall[,1])
```

```

tao0sq=100
v0=1
sigma0sq=1
# Initial sigma value for Gibbs sampling
n=dim(Rainfall)[1]
vn=v0+n
nDraws=5000

# Function for calculating tao_n which is used as argument for the std dev
for the normal distribution of mu
calcTaoN = function(sigmatq,tao0sq,n){
  return(1/(n/sigmatq+1/tao0sq))
}

calcMuN = function(sigmatq, tao0sq, mu0, mean, n) {
  w=(n/sigmatq)/(n/sigmatq+1/tao0sq)
  return(w*mean+(1-w)*mu0)
}

calcSigmaHat = function(v0, sigma0sq, data, mu, n) {
  return((v0*sigma0sq+sum((data-mu)^2))/(n+v0))
}

posteriorMatrix = matrix(0, nDraws, 2)
# Setting initial value of sigma^2 to 1
posteriorMatrix[1,2]=1
for (i in 1:nDraws) {
  posteriorMatrix[i,1] = rnorm(1, calcMuN(posteriorMatrix[i,2],tao0sq, mu0
, mean_rainfall, n),
                                calcTaoN(posteriorMatrix[i,2], tao0sq, n))

  if(i<nDraws) {
    drawX=rchisq(1,vn)
    posteriorMatrix[i+1,2]=vn*calcSigmaHat(v0, sigma0sq, Rainfall[,1], pos
teriorMatrix[i,1], n)/drawX
  }
}

# The posterior coverage
plot(posteriorMatrix[1001:nrow(posteriorMatrix),1], posteriorMatrix[1001:n
row(posteriorMatrix),2], xlab="Mu",
      ylab="Sigma^2")

## ii) AnalyzethedailyprecipitationusingyourGibbsamplerin(a)-i. Evaluate
the convergence of the Gibbs sampler
## by suitable graphical methods, for example by plotting the trajectories
of the sampled Markov chains.

iter=seq(1001,5000,1)
plot(iter, posteriorMatrix[1001:nrow(posteriorMatrix),1], type="l", xlab="
Iteration",
      ylab="Mu", main="Marginal posterior for mu")
plot(iter, posteriorMatrix[1001:nrow(posteriorMatrix),2], type="l", xlab="
Iteration",
      ylab="Sigma", main="Marginal posterior for sigma")

```



```
## b) Let us now instead assume that the daily precipitation {y1,...,yn} follow an iid two-component mixture
## of normals model:  $p(y_i \text{ given } \mu, \sigma^2, \pi) = \pi * N(y_i \text{ given } \mu_1, \sigma_1^2) + (1-\pi) * N(y_i \text{ given } \mu_2, \sigma_2^2)$ 
## where  $\mu = (\mu_1, \mu_2)$  and  $\sigma^2 = (\sigma_1^2, \sigma_2^2)$ 
## Use the Gibbs sampling data augmentation algorithm in NormalMixtureGibbs.R (available under Lecture 7 on the
## course page) to analyze the daily precipitation data. Set the prior hyperparameters suitably. Evaluate the
## convergence of the sampler.
```

```
# NormalMixtureGibbs.R with modifications
```

```
##### BEGIN USER INPUT #####
```

```
# Data options
```

```
x <- as.matrix(Rainfall[,1])
```

```
# Model options
```

```
nComp <- 2 # Number of mixture components
```

```
# Prior options
```

```
alpha <- rep(1,nComp) # Dirichlet(alpha)
```

```
# Obtained from Google, prior knowledge
```

```
muPrior <- c(14.79, 17.6) # Prior mean of mu
```

```
tau2Prior <- rep(100,nComp) # Prior std of mu
```

```
sigma2_0 <- rep(var(x),nComp) # s20 (best guess of sigma2)
```

```
nu0 <- rep(1,nComp) # degrees of freedom for prior on sigma2
```

```
# MCMC options
```

```
nIter <- 1000 # Number of Gibbs sampling draws
```

```
# Plotting options
```

```
plotFit <- TRUE
```

```
lineColors <- c("blue", "green", "magenta", 'yellow')
```

```
sleepTime <- 0.01 # Adding sleep time between iterations for plotting
```

```
##### END USER INPUT #####
```

```
##### Defining a function that simulates from the
```

```
rScaledInvChi2 <- function(n, df, scale){
```

```
  return((df*scale)/rchisq(n,df=df))
```

```
}
```

```
##### Defining a function that simulates from a Dirichlet distribution
```

```
rDirichlet <- function(param){
```

```
  nCat <- length(param)
```

```
  piDraws <- matrix(NA,nCat,1)
```

```
  for (j in 1:nCat){
```

```
    piDraws[j] <- rgamma(1,param[j],1)
```

```
  }
```

```
  piDraws = piDraws/sum(piDraws) # Dividing every column of piDraws by the sum of the elements in that column.
```

```
  return(piDraws)
```

```
}
```

```

# Simple function that converts between two different representations of the mixture allocation
S2alloc <- function(S){
  n <- dim(S)[1]
  alloc <- rep(0,n)
  for (i in 1:n){
    alloc[i] <- which(S[i,] == 1)
  }
  return(alloc)
}

# Initial value for the MCMC
nObs <- length(Rainfall[,1])
S <- t(rmultinom(nObs, size = 1, prob = rep(1/nComp,nComp))) # nObs-by-nComp matrix with component allocations.
mu <- quantile(x, probs = seq(0,1,length = nComp))
sigma2 <- rep(var(x),nComp)
probObsInComp <- rep(NA, nComp)

# Setting up the plot
xGrid <- seq(min(x)-1*apply(x,2,sd),max(x)+1*apply(x,2,sd),length = 100)
xGridMin <- min(xGrid)
xGridMax <- max(xGrid)
mixDensMean <- rep(0,length(xGrid))
effIterCount <- 0
ylim <- c(0,2*max(hist(x)$density))
param_matrix=matrix(0,4,nIter)
rownames(param_matrix)=c("Mu1", "Mu2", "Sigma1", "Sigma2")

for (k in 1:nIter){
  message(paste('Iteration number:',k))
  alloc <- S2alloc(S) # Just a function that converts between different representations of the group allocations
  nAlloc <- colSums(S)
  print(nAlloc)
  # Update components probabilities
  pi <- rDirichlet(alpha + nAlloc)

  # Update mu's
  for (j in 1:nComp){
    precPrior <- 1/tau2Prior[j]
    precData <- nAlloc[j]/sigma2[j]
    precPost <- precPrior + precData
    wPrior <- precPrior/precPost
    muPost <- wPrior*muPrior + (1-wPrior)*mean(x[alloc == j])
    tau2Post <- 1/precPost
    mu[j] <- rnorm(1, mean = muPost, sd = sqrt(tau2Post))
  }
  param_matrix[1,k]=mu[1]
  param_matrix[2,k]=mu[2]

  # Update sigma2's
  for (j in 1:nComp){
    sigma2[j] <- rScaledInvChi2(1, df = nu0[j] + nAlloc[j],

```

```

scale = (nu0[j]*sigma2_0[j] +
sum((x[alloc == j] - mu[j])^2))
/(nu0[j] + nAlloc[j]))
}
param_matrix[3,k]=sigma2[1]
param_matrix[4,k]=sigma2[2]

# Update allocation
for (i in 1:nObs){
  for (j in 1:nComp){
    probObsInComp[j] <- pi[j]*dnorm(x[i], mean = mu[j], sd = sqrt(sigma2
[j]))
  }
  S[i,] <- t(rmultinom(1, size = 1, prob = probObsInComp/sum(probObsInC
omp)))
}

# Printing the fitted density against data histogram
if (plotFit && (k%1 == 0)){
  effIterCount <- effIterCount + 1
  hist(x, breaks = 20, freq = FALSE, xlim = c(xGridMin,xGridMax), main =
paste("Iteration number",k),
  ylim = ylim)
  mixDens <- rep(0,length(xGrid))
  components <- c()
  for (j in 1:nComp){
    compDens <- dnorm(xGrid,mu[j],sd = sqrt(sigma2[j]))
    mixDens <- mixDens + pi[j]*compDens
    lines(xGrid, compDens, type = "l", lwd = 2, col = lineColors[j])
    components[j] <- paste("Component ",j)
  }
  mixDensMean <- ((effIterCount-1)*mixDensMean + mixDens)/effIterCount

  lines(xGrid, mixDens, type = "l", lty = 2, lwd = 3, col = 'red')
  legend("topright", box.lty = 1, legend = c("Data histogram",components
, 'Mixture'),
  col = c("black",lineColors[1:nComp], 'red'), lwd = 2)
}
}

hist(x, breaks = 20, freq = FALSE, xlim = c(xGridMin,xGridMax), main = "Fi
nal fitted density")
lines(xGrid, mixDensMean, type = "l", lwd = 2, lty = 4, col = "red")
lines(xGrid, dnorm(xGrid, mean = mean(x), sd = apply(x,2,sd)), type = "l",
lwd = 2, col = "blue")
legend("topright", box.lty = 1, legend = c("Data histogram","Mixture densi
ty","Normal density"),
col=c("black","red","blue"), lwd = 2)
plot(param_matrix[1,200:ncol(param_matrix)], type="l")
plot(param_matrix[2,200:ncol(param_matrix)], type="l")
plot(param_matrix[3,200:ncol(param_matrix)], type="l")
plot(param_matrix[4,200:ncol(param_matrix)], type="l")

```

```
## It seems like the sampler has converged towards a mixture distribution
which resembles the histogram of
## the data. The mode of the distribution is approximately at 20*1/100 inc
hes per day. The mixture density
## function seems to resemble the reality more accurately than the normal
density function. It seems reasonable
## to apply a mixture distribution to this type of data since rain is not
a constant occurrence but can happen
## on some days, and on some days not. When going through the iterations i
t is apparent that the mixture
## distribution converges quite quickly.
```

```
## c) Plot the following densities in one figure: 1) a histogram or kernel
density estimate of the data.
## 2) Normal density of  $N(y_i \text{ given } \mu \text{ and } \sigma^2)$  in a); 3) Mixture of no
rmal density
##  $p(y_i \text{ given } \mu, \sigma^2, \pi)$  in b). Base your plots on the mean over all
posterior draws.
```

```
hist(x, breaks = 20, freq = FALSE, xlim = c(xGridMin,xGridMax), main = "Fi
nal density")
lines(xGrid, mixDensMean, type = "l", lwd = 2, lty = 4, col = "red")
lines(xGrid, dnorm(xGrid, mean = mean(posteriorMatrix[,1]), sd = mean(sqrt
(posteriorMatrix[,2]))),
      type = "l", lwd = 2, col = "blue")
legend("topright", box.lty = 1, legend = c("Data histogram","Mixture densi
ty","Normal density"),
      col=c("black","red","blue"), lwd = 2)
```

```
## As seen in the new plot, where the only difference is the blue line, th
e resembles to the previous plot is
## obvious. The blue curve has not changed at all which is due to the fact
that the mean of Gibbs sampled data,
## when iterations go towards infinity, converges to the real mean of the
data.
```

Assignment 2 – Linear regression with Poisson, ML-estimator, optim with Poisson, simulating through RWM Hastings, myFunction(...), simulate from predictive distribution

```
## Assignment 2: Consider the following Poisson regression model  $y_i$  given
 $\beta \sim \text{Poisson}(\exp(x_i^T \beta))$ ,  $i=1, \dots, n$ 
## where  $y_i$  is the count for the  $i$ th observation in the sample and  $x_i$  is the  $p$ -dimen
sional vector with covariate observations
## for the  $i$ th observation. Use the data set eBayNumberOfBidderData.dat. T
his dataset contains observations from 1000
## eBay auctions of coins. The response variable is nBids and records the
number of bids in each auction. The
## remaining variables are features/covariates (x):
```

```
## a) Obtain the maximum Likelihood estimator of  $\beta$  in the Poisson regre
ssion model for the eBay data
## [Hint: glm.R, don't forget that glm() adds its own intercept so don't i
nput the covariate Const]. Which
## covariates are significant?
```

```

# Read data
ebay = read.table("ebayNumberOfBidderData.dat", header=TRUE)
data = ebay[, -2]

# Create model
model = glm(nBids~., family="poisson", data=data)
print(model$coefficients)
summary(model)

## The covariates that are significant are VerifyID, Sealed, MajBlem, LogBook, MinBidShare.

## b) Let's now do a Bayesian analysis of the Poisson regression. Let the
prior be  $\text{Beta} \sim N(0, 100 * (X^T X)^{-1})$  where  $X$ 
## is the  $n \times p$  covariate matrix. This is a commonly used prior which is called Zellner's g-prior. Assume first that
## the posterior density is approximately multivariate normal:  $\text{Beta} \text{ given } y \sim N(\text{Beta}^*, J_y(\text{Beta}^*)^{-1})$  where  $\text{Beta}^*$ 
## is the posterior mode and  $J_y(\text{Beta}^*)$  is the negative Hessian at the posterior mode.  $\text{Beta}^*$  and  $J$  can be obtained
## by numerical optimization (optim.R) exactly like you already did for the logistic regression in Lab2.

library(mvtnorm)
# Defining constants
X = as.matrix(ebay[, 2:ncol(ebay)])
Y = ebay[, 1]
nFeatures = dim(X)[2]
covNames = names(ebay[, 2:ncol(ebay)])

# Constructing prior
mu_prior = rep(0, nFeatures)
sigma_prior = 100 * solve(t(X) %*% X)

# Defining function for returning the log posterior
logPostPoisson = function(beta, Y, X, mu, sigma) {
  n = length(Y)
  XBeta = beta %*% t(X)
  # Defining loglikelihood
  logLike <- sum(-log(factorial(Y)) + XBeta * Y - exp(XBeta))
  # Defining prior
  prior = dmvnorm(beta, mean = mu, sigma = sigma, log = TRUE)
  # Adding Loglikelihood and Logprior together. Since it is log both of them are added instead of multiplied
  return(logLike + prior)
}

# Defining initial values to be passed on to the optimizer
set.seed(12345)
initVals = rnorm(dim(X)[2])

# Finding the optimized betavector
optimResult = optim(initVals, logPostPoisson, Y=Y, X=X, mu=mu_prior, sigma=

```

```
=sigma_prior, method=c("BFGS"),
                      control=list(fnscale=-1), hessian=TRUE)
```

*# Defining the values of interest*

```
postMode = optimResult$par
postCov = -solve(optimResult$hessian)
names(postMode) = covNames
approx_PostStd = sqrt(diag(postCov))
names(approx_PostStd) = covNames
print("The posterior mode is:")
print(postMode)
print("The approximated standard deviations are:")
print(approx_PostStd)
```

*## Through optimization we have obtained the optimal betavector as well as the hessian evaluated at the posterior mode.*

*## c) Now, let's simulate from the actual posterior of beta using the Metropolis algorithm and compare with the approximate results in b). Program a general function that uses the Metropolis algorithm to generate random draws from an arbitrary posterior density. In order to show that it is a general function for any model, I will denote the vector of model parameters by theta. Let the proposal density be the multivariate normal density mentioned in Lecture 8 (random walk Metropolis):  $\Theta_p$  given  $\Theta_{i-1} \sim N(\Theta_{i-1}, c \cdot \text{Cov})$  where  $\text{Cov} = \text{Jy}(\text{Beta})^{(-1)}$  obtained in b). The value c is a tuning parameter and should be an input to your Metropolis function. The user of your Metropolis function should be able to supply her own posterior density function, not necessarily for the Poisson regression, and still be able to use your Metropolis function. This is not so straightforward, unless you have come across function objects in R and the triple dot (...) wildcard argument. I have posted a note (HowToCodeRWM.pdf) on the course web page that describes how to do this in R. Now, use your new Metropolis function to sample from the posterior of beta in the Poisson regression for the eBay dataset. Assess MCMC convergence by graphical methods.*

*# Defining function for sampling through metropolishastings*

```
RVMSampler = function(previousVal, postCov, c, myFunction, ...) {
  proposalVal=rmvnorm(1, mean=previousVal, sigma=c*postCov)
  alpha=min(1, exp(myFunction(proposalVal,...)-myFunction(previousVal, ...)))
  u=runif(1)
  if(u < alpha) {
    return(proposalVal)
  } else {
    return(previousVal)
  }
}
```

```

nDraws=5000
beta_matrix = matrix(0, nDraws, ncol(X))
# Setting initial values of beta to same initVals as in the optimizer (taken randomly from normal distrib)
beta_matrix[1,]=initVals
c=0.5
set.seed(12345)

for(i in 1:nDraws) {
  if(i<nDraws) {
    beta_matrix[i+1,]=RVMSampler(beta_matrix[i,], postCov, c, logPostPoisson, Y, X, mu_prior, sigma_prior)
  }
}

iter=seq(1,nDraws,1)
par(mfrow=c(3,3))
for (i in 1:9) {
  plot(iter, beta_matrix[,i], type="l", main=paste("Convergence plot for covariate", covNames[i]),
        ylab=covNames[i])
}
par(mfrow=c(1,1), new=FALSE)

# Calculating distinct rows and dividing by total rows to get average acceptance probability
avg_alpha=dim(beta_matrix[!duplicated(beta_matrix),,])[1]/dim(beta_matrix)[1]

## As seen in the convergence plots the covariates oscillate around the same value which was found in the previous
## problem where the optimal beta values were found through optimization. Since the variable c should be chosen
## in a way to acquire an average acceptance rate of approximately 25-30%, the average acceptance rate were
## calculated to approximately 33 % which is deemed to be sufficiently satisfying.

## d) Use the MCMC draws from c) to simulate from the predictive distribution of the number of bidders in a new
## auction with the characteristics below. Plot the predictive distribution. What is the probability of no bidders
## in this new auction? Use vector x=c(1,1,1,1,0,0,0,1,0.5)

obs_X=c(1,1,1,1,0,0,0,1,0.5)
# Removing first 1000 rows since they are before the start of the convergence
approx_post_beta=beta_matrix[1001:nrow(beta_matrix),]
mean_vector=exp(approx_post_beta%%obs_X)
set.seed(12345)
pred_distrib_bidder=rpois(10000, mean_vector)
barplot(table(pred_distrib_bidder),
        main="Histogram of the predictive distribution of no. of bidders",

```



```
xlab="No. of bidders")
# Calculating the probability of no bidders with the given characteristics
prob_noBidders=sum(pred_distrib_bidder==0)/length(pred_distrib_bidder)
print(prob_noBidders)

## As seen in the predictive distribution the majority of cases given the
specified characteristics, will result in
## either 0 or 1 bidder with the probability decreasing for additional bid
ders. The calculated probability for
## no bidder is 0.3581.
```

## Lab 4 – HMC with Stan

Assignment 1 – AR(1)-process, stanmodel with AR(1)-process, Poisson conditioned on AR(1)-process in stan, plot with data & posterior mean & credible intervals over time

```
## Assignment 1:
## a) Write a function in R that simulate data from the AR(1)-process:  $x_t = \mu + \phi(x_{t-1} - \mu) + \epsilon(t)$ ,
##  $\epsilon(t) \sim N(0, \sigma^2)$ , for given values of  $\mu$ ,  $\phi$ , and  $\sigma^2$ . Start
the process at  $x_1 = \mu$  and then simulate
## values for  $x_t$  for  $t=2, 3, \dots, T$  and return the vector  $x_{1:T}$  containing all
time points. Use  $\mu=10$ ,  $\sigma^2=2$  and
##  $T=200$  and look at different realizations (simulation) of  $x_{1:T}$  for values
of  $\phi$  between  $-1$  and  $1$  (this is the
## interval of  $\phi$  where the AR-process is stable). Include a plot of at least
one realization in the report. What
## effect does the value of  $\phi$  have on  $x_{1:t}$ 

#install.packages("rstan")
mu=10
sigma_sq=2
T=200
x_init=mu
phi_vector=seq(-0.9,0.9,0.1)
results_matrix=matrix(0,200,length(phi_vector))
results_matrix[1,]=x_init
counter=1
set.seed(12345)

AR_process_function=function(mu, sigma_sq, T, phi) {
  x_init=mu
  result=rep(0,T)
  result[1]=x_init
  for (i in 2:T) {
    epsilon=rnorm(1,0,sqrt(sigma_sq))
    result[i]=mu+phi*(result[i-1]-mu)+epsilon
  }
  return(result)
}

results_matrix=matrix(0,T,length(phi_vector))
counter=1
for (phi in phi_vector) {
  results_matrix[,counter]=AR_process_function(mu,sigma_sq,T,phi)
}
```



```

    counter=counter+1
}
iter=seq(1,200,1)
counter=1
for (i in 1:length(phi_vector)) {
  if (counter %% 6 == 0) {
    plot(iter, results_matrix[,i], main="Plot of realization of AR-process",
        sub=paste("Phi =", phi_vector[i]),
        xlab="Iteration", ylab="Value", type="l", col="grey")
  }
  counter=counter+1
}

```

*## With phi-values below zero the process will oscillate faster but with phi-values above zero the process will be more correlated. The correlation between the different iterations increases as the phi-value becomes larger. This causes the oscillation to slow down and the process to move more slowly.*

*## b) Use your function from a) to simulate two AR(1)-processes,  $x_{1:T}$  with  $\phi=0.3$  and  $y_{1:T}$  with  $\phi=0.95$ . Now, treat the values of  $\mu$ ,  $\phi$  and  $\sigma^2$  as unknown and estimate them using MCMC. Implement Stan-code that samples from the posterior of the three parameters, using suitable non-informative priors of your choice. [Hint: Look at the time-series models examples in the Stan user's guide/reference manual, and note the different parametrizations used here.]*

*## i) Report the posterior mean, 95% credible intervals and the number of effective posterior samples for the three inferred parameters for each of the simulated AR(1)-process. Are you able to estimate the true values?*

*## ii) For each of the two data sets, evaluate the convergence of the samplers and plot the joint posterior of  $\mu$  and  $\phi$ . Comments?*

```

library(rstan)

x=rep(0,T)
y=rep(0,T)
set.seed(12345)
x=AR_process_function(mu, sigma_sq, T, 0.3)
set.seed(12345)
y=AR_process_function(mu, sigma_sq, T, 0.95)

StanModel= '
data {
  int<lower=0> N;
  vector[N] y;
}

parameters {

```

```

    real mu;
    real phi;
    real<lower=0> sigma;
  }
  model {
    for (n in 2:N)
      y[n] ~ normal(mu + phi * (y[n-1]-mu), sigma);
  }
  ;

data_x=list(N=T, y=x)
data_y=list(N=T, y=y)
fit_x=stan(model_code=StanModel, data=data_x)
fit_y=stan(model_code=StanModel, data=data_y)
postDraws_x <- extract(fit_x)
postDraws_y <- extract(fit_y)
print(fit_x)
print(fit_y)

# Do traceplots of the first chain
plot(postDraws_x$mu[1000:2000], postDraws_x$phi[1000:2000],ylab="phi", xlab="mu", main="Traceplot")

# Do traceplots of the first chain
plot(postDraws_y$mu[1000:2000],postDraws_y$phi[1000:2000],ylab="mu", xlab="mu",main="Traceplot")

## The posterior mean, number of effective samples as well as 95 % credible interval are shown above for both of the
## simulated AR(1)-processes. It is possible to estimate the true values of the parameters for the sample which
## used a phi=0.3 when obtaining the dataset used in the simulation. However, it is not as obvious to estimate
## the parameters' true values for the second sample where phi=0.95 were used to obtain the dataset used in this
## particular simulation. The credible intervals for the parameters in this simulation are very wide and it is
## difficult to predict with certainty the true value of the parameter. This might be due to the higher correlation
## between the lags caused by the higher value of phi.

## The convergence of the samplers are different. For the first sample which used phi=0.3, the convergence is
## evident whilst for the second sample the posterior distribution is not obvious. This correlates with the fact
## the credible intervals for the parameters on the second sample were very wide. What we can see from the
## posterior distribution obtained by the second sampler is that for lower values of phi the distribution centers
## around a value between 10 and 20. This is a behaviour similar to what is shown in the posterior for the first
## sampler, where phi was set to 0.3 initially, since this distribution was much tighter around the value of 10
## for mu.

```

## c) The data campy.dat contain the number of cases of campylobacter infections in the north of the province  
 ## Quebec (Canada) in four week intervals from January 1990 to the end of October 2000. It has 13 observations per  
 ## year and 140 observations in total. Assume that the number of infections  $c_t$  at each time point follows an  
 ## independent Poisson distribution when conditioned on a latent AR(1)-process  $x_t$ , that is  
 ##  $c_t$  given  $x_t \sim \text{Poisson}(\exp(x_t))$ , where  $x_t$  is an AR(1)-process as in a). Implement and estimate the model in Stan,  
 ## using suitable priors of your choice. Produce a plot that contains both the data and the posterior mean and  
 ## 95 % credible intervals for the latent intensity  $\theta_t = \exp(x_t)$  over time.  
 ## [Hint: Should  $x_t$  be seen as data or parameters]

```
campy=read.table("campy.dat", header=TRUE)
library(rstan)

StanModel_Pois = '
data {
  int<lower=0> T;
  int c[T];
}

parameters {
  real mu;
  real phi;
  real<lower=0> sigma;
  vector[T] x;
}

model {
  // Prior
  phi ~ uniform(-1,1);
  for (n in 2:T)
    x[n] ~ normal(mu + phi * (x[n-1]-mu), sigma);

  // Model/likelihood
  for (n in 1:T)
    c[n] ~ poisson(exp(x[n]));
}

generated quantities {
  vector[T] post_mean;
  post_mean = exp(x);
}
'

data=list(T=dim(campy)[1], c=campy$c)
fit_pois=stan(model_code=StanModel_Pois, data=data)
```

```

print(fit_pois)
pois_mean_list=fit_pois@.MISC$summary$msd
post_mean=pois_mean_list[grep("post_mean", rownames(pois_mean_list)),]

plot(campy$c, col="blue", ylab="No. of infected", xlab="Time")
points(post_mean[,1], col="black", type="l")

quantiles=fit_pois@.MISC$summary$quan
quantiles_post_mean=quantiles[grep("post_mean", rownames(quantiles)),]
cred_interval_post_mean=matrix(0,dim(quantiles_post_mean)[1], 2)
cred_interval_post_mean[:,1]=quantiles_post_mean[:,1]
cred_interval_post_mean[:,2]=quantiles_post_mean[:,ncol(quantiles_post_mean)]

lines(cred_interval_post_mean[:,1], col="gray", lty=1)
lines(cred_interval_post_mean[:,2], col="gray", lty=1)
title(main="Plot of data vs approximated posterior")
legend("topleft", box.lty= 1, pch=c(1,NaN,NaN), legend=c("Data", "Posterior mean", "95 % cred. interval"),
       col=c("blue", "black", "gray"), lwd=c(NaN,1,1), lty=c(NaN, 1, 1))

## As seen in the plot above the posterior mean follows the data accurately. Almost all of the datapoints are
## inside the credible intervals which aren't that wide which indicates that the approximated posterior
## resembles the reality shown by the data well.

## d) Now, assume that we have a prior belief that the true underlying intensity  $\theta_t$  varies more smoothly than
## the data suggests. Change the prior for  $\sigma^2$  so that it becomes informative about that the AR(1)-process
## increments  $\epsilon_t$  should be small. Re-estimate the model using Stan with the new prior and produce the same
## plot as in c). Has the posterior for  $\theta_t$  changed?

StanModel_Pois_Prior = '
data {
  int<lower=0> T;
  int c[T];
}

parameters {
  real mu;
  real phi;
  real<lower=0> sigma;
  vector[T] x;
}

model {
  // Prior
  phi ~ uniform(-1,1);
  sigma ~ scaled_inv_chi_square(140, 0.15);
  for (n in 2:T)
    x[n] ~ normal(mu + phi * (x[n-1]-mu), sigma);

```

```

// Model/likelihood
for (n in 1:T)
  c[n] ~ poisson(exp(x[n]));
}

generated quantities {
  vector[T] post_mean;
  post_mean = exp(x);
}

```

```

fit_pois_prior=stan(model_code=StanModel_Pois_Prior, data=data)
print(fit_pois_prior)
pois_mean_list_prior=fit_pois_prior@.MISC$summary$msd
post_mean_prior=pois_mean_list_prior[grep("post_mean", rownames(pois_mean_
list)),]

plot(campy$c, col="blue", ylab="No. of infected", xlab="Time")
points(post_mean_prior[,1], col="black", type="l")

quantiles_prior=fit_pois_prior@.MISC$summary$quan
quantiles_post_mean_prior=quantiles_prior[grep("post_mean", rownames(quant
iles)),]
cred_interval_post_mean_prior=matrix(0,dim(quantiles_post_mean)[1], 2)
cred_interval_post_mean_prior[,1]=quantiles_post_mean_prior[,1]
cred_interval_post_mean_prior[,2]=quantiles_post_mean_prior[,ncol(quantile
s_post_mean)]

lines(cred_interval_post_mean_prior[,1], col="gray", lty=1)
lines(cred_interval_post_mean_prior[,2], col="gray", lty=1)
title(main="Plot of data vs approximated posterior")
legend("topleft", box.lty= 1, pch=c(1,NaN,NaN), legend=c("Data", "Posterior mean", "95 % cred. interval"),
      col=c("blue", "black", "gray"), lwd=c(NaN,1,1), lty=c(NaN, 1, 1))

```

*## Now when we have specified a small prior for sigma it is notable in the new plot that the posterior mean varies less and moves more smoothly. The consequence of this is that more datapoints lie outside of the credible interval which suggests that the approximated posterior does not resemble the reality described by the data as accurately as before. However, by doing this one can avoid overfitting when the model is applied to a new dataset.*

## Exams

2017-05-30

Assignment 1 – Plot posterior (rice function), normal approx. through optim, simulation for new obs

*## a) Plot the posterior distribution of theta*

```
riceData <- c(1.556, 1.861, 3.135, 1.311, 1.877, 0.622, 3.219, 0.768, 2.358, 2.056)
```

*# Random number generator for the Rice distribution*

```
rRice <- function(n = 1, theta = 1, psi = 1){  
  x <- rnorm(n = n, mean = 0, sd = sqrt(psi))  
  y <- rnorm(n = n, mean = theta, sd = sqrt(psi))  
  return(sqrt(x^2+y^2))  
}
```

*# Function for calculating the log posterior distrib with theta prior set to 1*

```
logPosterior = function(data, theta, psi) {  
  bessel_factor=1  
  for (i in data) {  
    bessel_factor=bessel_factor*besselI(i*theta/psi, nu=0)  
  }  
  post=-log(psi)-1/(2*psi)*sum(data^2+theta^2)+log(bessel_factor)  
  return(post+0) # If prior is assumed to be constant we set the prior to 1 which in log scale yields 0  
}
```

```
gridWidth=0.01
```

```
theta_grid=seq(0,3,gridWidth)
```

```
posterior_distrib_log=sapply(theta_grid, logPosterior, data=riceData, psi=1)
```

```
posterior_distrib_norm=1/gridWidth*exp(posterior_distrib_log)/sum(exp(posterior_distrib_log))
```

```
sum(posterior_distrib_norm)
```

```
plot(theta_grid, posterior_distrib_norm, xlab=expression(theta), ylab="Density",  
      main="Posterior density of theta",  
      type="l", lwd=2)
```

*## b) Use numerical optimization to obtain a normal approx. of the posterior distrib of theta. Overlay curve*

*## from a) with the approximated normal distribution*

*# Defining initial values to be passed on to the optimizer*

```
set.seed(12345)
```

```
initVal = rnorm(1, mean=0, sd=1)
```

*# Finding the optimized betavector*

```
optimResult = optim(initVal, logPosterior, data=riceData, psi=1, method=c("L-BFGS-B"),  
                    control=list(fnscale=-1), lower=0, hessian=TRUE)
```

```

# Defining the values of interest
postMode = optimResult$par
postCov = as.numeric(-solve(optimResult$hessian))
print("The posterior mode is:")
print(postMode)
print("The approximated standard deviation is:")
print(postCov)
lines(theta_grid, dnorm(theta_grid, mean=postMode, sd=sqrt(postCov)), col=
"red", lwd=2)
legend(x = 1.8, y = 1, legend = c("True posterior", "Approximate posterior
"),
      col = c("black", "red"), lty = c(1,1), lwd = c(2,2), cex = 0.8)

## Answer: Not perfect approx but fairly good.

## c) Simulate distrib for new observation using normal approx in b)

nDraws=5000
set.seed(12345)
theta=rnorm(nDraws, mean=postMode, sd=sqrt(postCov))
pred_distrib=c()
for (i in theta) {
  pred_distrib=c(pred_distrib, rRice(theta=i))
}

hist(pred_distrib, breaks=100, xlab="Index", main="Predictive density of n
ew obs")

```

Assignment 2 – Model posterior data, plot posterior and compare with data, Gibbs for mixture of models, graphical methods for evaluating

*## a) Model posterior data with prior Gamma and Likelihood Poisson, plot the posterior*

*# We know that posterior mapping with gamma prior and poisson likelihood is gamma distributed*

```

sumBids=sum(bids)
n=length(bids)
alpha=1
beta=1
posterior_theta=dgamma(seq(3,4,0.001), alpha+sumBids, beta+n)
plot(seq(3,4,0.001), posterior_theta, type="l", lwd=2)

```

*# b) Investigate through graphical methods if Poisson model describes data well*

```

xGrid=seq(min(bids), max(bids))
data_norm=bidsCounts/sum(bidsCounts)
nDraws=5000
thetaDraws=rgamma(nDraws, alpha+sumBids, beta+n)
poissonDensity=rep(0, length(xGrid))
for (i in thetaDraws) {
  poissonDensity=poissonDensity+dpois(xGrid, lambda=i)
}

```

```

}

avgPoissonDensity=poissonDensity/nDraws
plot(xGrid, data_norm, xlab="No. of bids", ylab="Density", main="Fitted models", type="o", cex=0.8,
      ylim=c(0,0.25), lwd=2)
lines(xGrid, avgPoissonDensity, col="red", lwd=2, type="o")
legend(x=7, y=0.2, col=c("black", "red"), legend=c("Data", "Poisson mean density"), lty=c(1,1),
       lwd=c(2,2), pch=c("o", "o"))

## Terrible fit which the plot shows

## c) Use GibbsMixPois.R. Estimate the mixture of Poissons both with K=2 and K=3. nIter=5000.

GibbsMixPois <- function(x, nComp, alpha, alphaGamma, betaGamma, xGrid, nIter){

  # Gibbs sampling for a mixture of Poissons
  # Author: Mattias Villani, IDA, Linkoping University. http://mattiasvillani.com
  #
  # INPUTS:
  # x - vector with data observations (counts)
  # nComp - Number of mixture components to be fitted
  # alpha - The prior on the mixture component weights is  $w \sim \text{Dir}(\alpha, \alpha, \dots, \alpha)$ 
  # alphaGamma and betaGamma -
  # The prior on the mean (theta) of the Poisson mixture components is
  #  $\theta \sim \text{Gamma}(\alpha\text{Gamma}, \beta\text{Gamma})$  [rate parametrization of the Gamma dist]
  # xGrid - the grid of data values over which the mixture is evaluated and plotted
  # nIter - Number of Gibbs iterations
  #
  # OUTPUTS:
  # results$wSample - Gibbs sample of mixture component weights. nIter-by-nComp matrix
  # results$thetaSample - Gibbs sample of mixture component means. nIter-by-nComp matrix
  # results$mixDensMean - Posterior mean of the estimated mixture density over xGrid.

  ##### Defining a function that simulates from a Dirichlet distribution
  rDirichlet <- function(param){
    nCat <- length(param)
    thetaDraws <- matrix(NA, nCat, 1)
    for (j in 1:nCat){
      thetaDraws[j] <- rgamma(1, param[j], 1)
    }
    thetaDraws = thetaDraws/sum(thetaDraws) # Diving every column of Theta
  }

```



*Draws by the sum of the elements in that column.*

```
    return(thetaDraws)
  }
```

*# Simple function that converts between two different representations of the mixture allocation*

```
S2alloc <- function(S){
  n <- dim(S)[1]
  alloc <- rep(0,n)
  for (i in 1:n){
    alloc[i] <- which(S[i,] == 1)
  }
  return(alloc)
}
```

*# Initial values for the Gibbs sampling*

```
nObs <- length(x)
S <- t(rmultinom(nObs, size = 1, prob = rep(1/nComp,nComp))) # nObs-by-
nComp matrix with component allocations.
```

*theta <- rep(mean(x), nComp) # Each component is initialized at the mean of the data*

*# Setting up the grid where the mixture density is evaluated.*

```
mixDensMean <- rep(0,length(xGrid))
effIterCount <- 0
```

*# Setting up matrices to store the draws*

```
wSample <- matrix(0, nIter, nComp)
thetaSample <- matrix(0, nIter, nComp)
probObsInComp <- rep(NA, nComp)
```

*# Setting up the priors - the same prior for all components*

```
alpha <- rep(alpha, nComp)
alphaGamma <- rep(alphaGamma, nComp)
betaGamma <- rep(betaGamma, nComp)
```

*# HERE STARTS THE ACTUAL GIBBS SAMPLING*

```
for (k in 1:nIter){
  message(paste('Iteration number:',k))
  alloc <- S2alloc(S) # Function that converts between different represe
ntations of the group allocations
  nAlloc <- colSums(S)
```

*# Step 1 - Update components probabilities*

```
w <- rDirichlet(alpha + nAlloc)
wSample[k,] <- w
```

*# Step 2 - Update theta's in Poisson components*

```
for (j in 1:nComp){
  theta[j] <- rgamma(1, shape = alphaGamma + sum(x[alloc == j]), rate
= betaGamma + nAlloc[j])
}
thetaSample[k,] <- theta
```

```

# Step 3 - Update allocation
for (i in 1:nObs){
  for (j in 1:nComp){
    probObsInComp[j] <- w[j]*dpois(x[i], lambda = theta[j])
  }
  S[i,] <- t(rmultinom(1, size = 1, prob = probObsInComp/sum(probObsInComp)))
}

# Computing the mixture density at the current parameters, and averaging that over draws.
effIterCount <- effIterCount + 1
mixDens <- rep(0,length(xGrid))
for (j in 1:nComp){
  compDens <- dpois(xGrid, lambda = theta[j])
  mixDens <- mixDens + w[j]*compDens
}
mixDensMean <- ((effIterCount-1)*mixDensMean + mixDens)/effIterCount
}
return(results = list(wSample = wSample, thetaSample = thetaSample, mixDensMean = mixDensMean))
}

```

```

result_comp2=GibbsMixPois(bids, nComp=2, alpha=1, alphaGamma = alpha, betaGamma = beta,
                           xGrid=xGrid, nIter=500)
result_comp3=GibbsMixPois(bids, nComp=3, alpha=1, alphaGamma = alpha, betaGamma = beta,
                           xGrid=xGrid, nIter=500)

```

## c) Use graphical methods to investigate if mixture of poissons fits data well. Is  $K=2$  enough or should we use  $K=3$ ?

```

plot(xGrid, data_norm, xlab="No. of bids", ylab="Density", main="Fitted models", type="o",
      ylim=c(0,0.25), lwd=2)
lines(xGrid, result_comp2$mixDensMean, col="red", lwd=2, type="o")
lines(xGrid, result_comp3$mixDensMean, col="gray", lwd=2, type="o")
legend(x=7, y=0.2, col=c("black", "red", "gray"),
       legend=c("Data", "Mixture density with 2 components", "Mixture density with 3 components"),
       lty=c(1,1,1), lwd=c(2,2, 2), pch=c("o", "o", "o"), cex=1)

```

## Good enough with 2 components in the mixture density

Assignment 3 – Linear regression (cars), marginal distributions, interpretation of cred. Interval, predictive simulation

```

#####
##### Problem 3 #####
#####

```

```

# Reading the cars data from file
load("cars.RData")

library(mvtnorm)

# Defining a function that simulates from the scaled inverse Chi-square distribution
rScaledInvChi2 <- function(n, df, scale){
  return((df*scale)/rchisq(n,df=df))
}

BayesLinReg <- function(y, X, mu_0, Omega_0, v_0, sigma2_0, nIter){
  # Direct sampling from a Gaussian linear regression with conjugate prior
  :
  #
  #  $\beta \mid \sigma^2 \sim N(\mu_0, \sigma^2 \text{inv}(\Omega_0))$ 
  #  $\sigma^2 \sim \text{Inv-}\chi^2(v_0, \sigma^2_0)$ 
  #
  # Author: Mattias Villani, IDA, Linköping University. http://mattiasvillani.com
  #
  # INPUTS:
  # y - n-by-1 vector with response data observations
  # X - n-by-nCovs matrix with covariates, first column should be ones if you want an intercept.
  # mu_0 - prior mean for beta
  # Omega_0 - prior precision matrix for beta
  # v_0 - degrees of freedom in the prior for sigma2
  # sigma2_0 - location ("best guess") in the prior for sigma2
  # nIter - Number of samples from the posterior (iterations)
  #
  # OUTPUTS:
  # results$betaSample - Posterior sample of beta. nIter-by-nCovs matrix
  # results$sigma2Sample - Posterior sample of sigma2. nIter-by-1 vector

  # Compute posterior hyperparameters
  n = length(y) # Number of observations
  nCovs = dim(X)[2] # Number of covariates
  XX = t(X)%*%X
  betaHat <- solve(XX, t(X)%*%y)
  Omega_n = XX + Omega_0
  mu_n = solve(Omega_n, XX%*%betaHat + Omega_0%*%mu_0)
  v_n = v_0 + n
  sigma2_n = as.numeric((v_0*sigma2_0 + ( t(y)%*%y + t(mu_0)%*%Omega_0%*%mu_0 - t(mu_n)%*%Omega_n%*%mu_n))/v_n)
  invOmega_n = solve(Omega_n)

  # The actual sampling
  sigma2Sample = rep(NA, nIter)
  betaSample = matrix(NA, nIter, nCovs)
  for (i in 1:nIter){

```

```

# Simulate from  $p(\sigma^2 \mid y, X)$ 
sigma2 = rScaledInvChi2(n=1, df = v_n, scale = sigma2_n)
sigma2Sample[i] = sigma2

# Simulate from  $p(\beta \mid \sigma^2, y, X)$ 
beta_ = rmvnorm(n=1, mean = mu_n, sigma = sigma2*invOmega_n)
betaSample[i,] = beta_

}
return(results = list(sigma2Sample = sigma2Sample, betaSample=betaSample
))
}

## a) Linear regression problem with given dataset. Use Mattias function to
derive joint posterior.
## i) Plot marginal distributions of each param
## ii) Compute point estimates for each regression coefficient assuming Loss
function
## iii) Construct 95 % equal tail probability intervals for each parameter
and interpret them

y=cars$mpg
x=as.matrix(cars[2:ncol(cars)])
mu_0=c(0,0,0,0)
omega_0=0.01*diag(x=4)
nu_0=1
sigma_sq_0=36
jointPostDistrib=BayesLinReg(y, x, mu_0, omega_0, nu_0, sigma_sq_0, 1000)
hist(jointPostDistrib$sigma2Sample, breaks=10, main=paste("Marginal distribution of",
expression(sigma^2)),
xlab=expression(sigma^2))
par(mfrow=c(2,2))
for(i in 1:4) {
hist(jointPostDistrib$betaSample[,i], breaks=10, main=paste("Marginal distribution of ",
expression(beta), i,
sep=""), xlab=
paste(expression(beta),i, sep=""))
}
title("Marginal distributions of the different betavalues", line=-1, outer=TRUE)
par(mfrow=c(1,1))

# Linear Loss function is posterior median
median(jointPostDistrib$sigma2Sample)
median(jointPostDistrib$betaSample[,1])
median(jointPostDistrib$betaSample[,2])
median(jointPostDistrib$betaSample[,3])
median(jointPostDistrib$betaSample[,4])

# Prediction intervals for each param
quantile(jointPostDistrib$sigma2Sample, c(0.025, 0.975))
quantile(jointPostDistrib$betaSample[,1], c(0.025, 0.975))
quantile(jointPostDistrib$betaSample[,2], c(0.025, 0.975))

```

```

quantile(jointPostDistrib$betaSample[,3], c(0.025, 0.975))
quantile(jointPostDistrib$betaSample[,4], c(0.025, 0.975))

## Answer: Interpretation of the credible interval for weight [-4.759964,
-1.531457]. A one unit increase of weight
## lowers the amount of miles per gallon between -4.759964 and -1.531457 w
ith 95 % posterior probability.

## b) Investigate if effect on mpg is different in cars with six cylinders
compared to cars with 8 cylinders

hist(jointPostDistrib$betaSample[,4]-jointPostDistrib$betaSample[,3], 50)
quantile(jointPostDistrib$betaSample[,4]-jointPostDistrib$betaSample[,3],
c(0.025, 0.975))

## Answer: Since 0 is present in interval we can not say that there is a d
ifference between 8 and 6 cylinders
## with 95 % posterior probability.

## c) Compute by simulation predictive distrib for a new car 4 cylinders a
nd weight=3.5

new_x=c(1,3.5,0,0)
pred_y=rep(0,nIter)
for (i in 1:nIter) {
  pred_y[i]=sum(new_x*jointPostDistrib$betaSample[i,])+rnorm(1,sd=sqrt(joi
ntPostDistrib$sigma2Sample[i]))
}
hist(pred_y, breaks=40, freq=FALSE)

```

#### Assignment 4 – Maximizing posterior expected utility

```

## Maximizing posterior expected utility.

post_dens = function(x) {
  return(gamma(6+x)/gamma(13+x))
}

barplot(post_dens(seq(0,10,1)), type="l")

# x6=10 seems to yield a low enough probability to be an upper bound for s
um

posterior_prob=post_dens(seq(0,10))
posterior_prob=posterior_prob/sum(posterior_prob)
exp_util=c()
for (k in 0:10) {
  exp_util=c(exp_util,(2^k-3))
}

exp_post_dens=sum(posterior_prob*exp_util)
print(exp_post_dens)

```

2017-08-16

Assignment 1 – Plot posterior density (Cauchy), normal approx. through optimization, marginal posterior

*## a) Plot posterior density of theta, with normal prior and cauchy distribution as likelihood*

*# Reading the data vector yVect from file*

```
load(file = 'CauchyData.RData')
```

```
cauchydata=yVect
```

```
dCauchy <- function(x, theta = 0, gamma = 1){  
  return(dens = (1/(pi*gamma))*(1/(1+((x-theta)/gamma)^2)))  
}
```

```
dlognormal <- function(x, mu, sigma2){  
  return(dens = (1/(sqrt(2*pi*sigma2)*x))*exp((-1/(2*sigma2))*(log(x)-mu)^2))  
}
```

```
logPrior_theta = function(theta, mu, sigma_sq) {  
  return(dnorm(theta, mean=mu, sd=sqrt(sigma_sq), log=TRUE))  
}
```

```
logPosterior = function(data, mu, sigma_sq, theta=0, gamma=1) {  
  prior=logPrior(theta, mu, sigma_sq)  
  likelihood=dCauchy(data, theta, gamma)  
  likelihood=sum(log(likelihood))  
  return(likelihood + prior)  
}
```

```
mu=0
```

```
sigma_sq=100
```

```
gamma=1
```

```
gridWidth=0.01
```

```
theta_grid=seq(0,8,gridWidth)
```

```
posterior_distrib=sapply(theta_grid, logPosterior, data=cauchydata, mu=mu,  
sigma_sq=sigma_sq, gamma=1)
```

```
posterior_distrib=1/gridWidth*exp(posterior_distrib)/sum(exp(posterior_distrib))
```

```
plot(theta_grid, posterior_distrib, type="l", lwd=2, main="Posterior density for theta",  
xlab=expression(theta),  
ylab="Density")
```

*## b) gamma is unknown with prior Lognormal.*

```
set.seed(12345)
```

```
initVal = c(0,0)
```

```
logJointPosterior = function(joint, data, mu, sigma_sq) {  
  prior_theta=logPrior(joint[1], mu, sigma_sq)  
  prior_gamma=log(dlognormal(joint[2], mu, sigma_sq))  
  likelihood=dCauchy(data, joint[1], joint[2])  
  likelihood=sum(log(likelihood))  
}
```

```

    return(likelihood + prior_theta + prior_gamma)
}

# Finding the optimized theta and gamma
optimResult = optim(initVal, logJointPosterior, data=cauchydata, mu=mu, sigma_sq=sigma_sq, method=c("L-BFGS-B"),
                    control=list(fnscale=-1), lower=c(-Inf, 0.001), upper=c(Inf, Inf), hessian=TRUE)

# Defining the values of interest
postMode = optimResult$par
postCov = -solve(optimResult$hessian)
names(postMode)=c("Theta", "Gamma")
print("The posterior mode is:")
print(postMode)
print("The approximated standard deviation is:")
print(postCov)

## c) Use normal approx in 1b) to obtain marginal posterior for the 99 % percentile of the cauchy distrib
## theta + gamma * tan(pi*(0.99-0.5))

library(rmvnorm)
normal_approx=rmvnorm(5000, mean=postMode, sigma=postCov)
cauchy_distrib=normal_approx[,1]+normal_approx[,2]*tan(pi*(0.99-0.5))
hist(cauchy_distrib, breaks=50, main="Marginal distribution of special case of cauchy", xlab="Function value")

```

Assignment 2 – Bayes Linear Regression, decision with loss function, predictive distribution

```

# Reading the data from file
library(MASS)
BostonHousing = Boston
y = BostonHousing$medv
X = cbind(1,BostonHousing[,1:13]) # Adding a column of ones for the intercept
names(X)[1] <- "intercept"
covNames <- names(X)
y <- as.numeric(y)
X <- as.matrix(X)

library(mvtnorm)

# Defining a function that simulates from the scaled inverse Chi-square distribution
rScaledInvChi2 <- function(n, df, scale){
  return((df*scale)/rchisq(n,df=df))
}

BayesLinReg <- function(y, X, mu_0, Omega_0, v_0, sigma2_0, nIter){
  # Direct sampling from a Gaussian linear regression with conjugate prior
  :
  #
  # beta | sigma2 ~ N(mu_0, sigma2*inv(Omega_0))
  # sigma2 ~ Inv-Chi2(v_0,sigma2_0)

```

```

#
# Author: Mattias Villani, IDA, Linköping University. http://mattiasvillani.com
#
# INPUTS:
# y - n-by-1 vector with response data observations
# X - n-by-nCovs matrix with covariates, first column should be ones if you want an intercept.
# mu_0 - prior mean for beta
# Omega_0 - prior precision matrix for beta
# v_0 - degrees of freedom in the prior for sigma2
# sigma2_0 - location ("best guess") in the prior for sigma2
# nIter - Number of samples from the posterior (iterations)
#
# OUTPUTS:
# results$betaSample - Posterior sample of beta. nIter-by-nCovs matrix
# results$sigma2Sample - Posterior sample of sigma2. nIter-by-1 vector

# Compute posterior hyperparameters
n = length(y) # Number of observations
nCovs = dim(X)[2] # Number of covariates
XX = t(X)%*%X
betaHat <- solve(XX,t(X)%*%y)
Omega_n = XX + Omega_0
mu_n = solve(Omega_n,XX%*%betaHat+Omega_0%*%mu_0)
v_n = v_0 + n
sigma2_n = as.numeric((v_0*sigma2_0 + ( t(y)%*%y + t(mu_0)%*%Omega_0%*%mu_0 - t(mu_n)%*%Omega_n%*%mu_n))/v_n)
invOmega_n = solve(Omega_n)

# The actual sampling
sigma2Sample = rep(NA, nIter)
betaSample = matrix(NA, nIter, nCovs)
for (i in 1:nIter){

  # Simulate from p(sigma2 | y, X)
  sigma2 = rScaledInvChi2(n=1, df = v_n, scale = sigma2_n)
  sigma2Sample[i] = sigma2

  # Simulate from p(beta | sigma2, y, X)
  beta_ = rmvnorm(n=1, mean = mu_n, sigma = sigma2*invOmega_n)
  betaSample[i,] = beta_

}
return(results = list(sigma2Sample = sigma2Sample, betaSample=betaSample))
}

mu_0=rep(0, ncol(X))
omega_0=0.01*diag(ncol(X))
v_0=1
sigma2_0=36

```



```

nIter=5000

bayes_lin_results=BayesLinReg(y, X, mu_0, omega_0, v_0, sigma2_0, nIter)
# Under quadratic loss, posterior mean is point estimate
beta_estimates=rep(0,ncol(X))
beta_credIntervals=matrix(0, ncol(X), 2)
for (i in 1:ncol(X)) {
  beta_estimates[i]=mean(bayes_lin_results$betaSample[,i])
  beta_credIntervals[i,]=quantile(bayes_lin_results$betaSample[,i], c(0.025, 0.975))
}
sigma_estimate=mean(bayes_lin_results$sigma2Sample)
sigma_credInterval=quantile(bayes_lin_results$sigma2Sample, c(0.025, 0.975))
rownames(beta_credIntervals)=covNames
beta_credIntervals[which(rownames(beta_credIntervals)=="rm"), ]

## Interpretation: for one unit increase of rooms the housing prices will rise between 3991,475 and 5009,826 dollars
## with 95 % posterior probability.

## b) Owner of house 381 is considering selling their house. Bought house for 10400

old_obs=as.vector(X[381,])
new_obs=old_obs
new_obs[2]=10
pred_draw=rep(0,nIter)
for (i in 1:nIter) {
  pred_draw[i]=bayes_lin_results$betaSample[i,]%*%new_obs+rnorm(1, mean=0, sd=sqrt(bayes_lin_results$sigma2Sample[i]))
}
pred_mean=mean(pred_draw)
hist(pred_draw, breaks=50)
quantile(pred_draw, c(0.025, 0.975))
sum(pred_draw>=30)/nIter

## c) See paper.

```

Assignment 4 – Simulation from posterior (Poisson with Gamma prior), simulation predictive distribution, maximizing posterior utility

## a) Simulate 1000 draws from the posterior distrib of theta using conjugate prior for theta with mean 250 and std = 50. Poisson Likelihood.

```

data=c(220,323,174,229)
alpha=25
beta=0.1
n=length(data)

logPriorGamma = function(theta, alpha, beta) {
  return(dgamma(theta, 50, beta, log=TRUE))
}

```

```

}

logLike = function(data, theta) {
  n=length(data)
  first=sum(data)*log(theta)
  second=theta*n
  third=0
  for (i in data) {
    for (j in 1:i) {
      third=third+log(j)
    }
  }
  return(first-second-third)
}

logPosterior = function(data, theta, alpha, beta) {
  prior=logPriorGamma(theta, alpha, beta)
  likelihood=logLike(data, theta)
  return(likelihood + prior)
}

# Conjugate prior for poisson is Gamma(alpha, beta), we know that posterior is Gamma(alpha + sum(data), beta+n)
post_draws=rgamma(1000, alpha+sum(data), beta+n)
hist(post_draws, main="Posterior distribution of theta", xlab=expression(theta))

## b) Simulate 1000 draws from the predictive distrib of next quarter's demand, X5, and plot the draws
## in histogram.

q5=rpois(1000, post_draws)
hist(q5, breaks=50, main="Predictive distribution of quarter 5", xlab="Qty")
sum(q5<=200)/1000

## c)

utility <- function(a,X5){
  util = rep(0,length(X5))
  util[X5<=a] = 10*X5[X5<=a]-(a-X5[X5<=a])
  util[X5>a] = 10*a-0.05*(X5[X5>a]-a)^2
  return(util)
}

mean(q5)
a=seq(136,336,1)
results = matrix(0,length(q5),length(a))
count=1
nameVec=rep(0,length(a))
for (i in a) {
  results[,count]=utility(i,q5)
  nameVec[count]=as.character(i)
}

```

```

    count=count+1
}
opt_vector=matrix(0,1,length(a))
for (i in 1:length(a)) {
  opt_vector[i]=mean(results[,i])
}
colnames(opt_vector)=nameVec
opt_decision=as.numeric(opt_vector[,which(opt_vector==max(opt_vector))])
names(opt_vector[,which(opt_vector==max(opt_vector))])
plot(a, opt_vector, type="l", lwd=1, col="red")
abline(v=as.numeric(names(opt_vector[,which(opt_vector==max(opt_vector))])),
  col="blue")

```

2017-10-27

Assignment 1 – Plot posterior distrib (beta symmetric prior, exponential likelihood), joint posterior distribution through optim, how to choose model?

*## a) Likelihood: Beta symmetric, prior, expon(1). Plot posterior distrib.*

```

thetaGrid=seq(0.01, 15, length=1000)
data=yProp
lambda=1

logPriorExp = function(theta, lambda) {
  return(dexp(theta, rate=lambda, log=TRUE))
}

logPosterior = function(x, theta, lambda) {
  prior=logPriorExp(theta, lambda)
  likelihood=sum(dbeta(x, theta, theta, log=TRUE))
  return(likelihood+prior)
}

theta_post=apply(thetaGrid, logPosterior, x=data, lambda=lambda)
theta_post_norm=1/((15-0.01)/1000)*exp(theta_post)/sum(exp(theta_post))
plot(thetaGrid, theta_post_norm, type="l", lwd=2, xlab=expression(theta),
  ylab="Posterior density")

# Zero to 1 loss means posterior mode is the optimal point estimator

index=which(theta_post_norm==max(theta_post_norm))
opt_theta=thetaGrid[index]
print(opt_theta)

```

*## Optimal theta is around 4.481491*

*## b) Theta1 and theta2 are independent apriori. Plot joint posterior distribution*

```

logPosteriorMult = function(theta, x, lambda) {
  theta1=theta[1]
  theta2=theta[2]
  prior1=logPriorExp(theta1, lambda)
  prior2=logPriorExp(theta2, lambda)
}

```

```

    likelihood=sum(dbeta(x, theta1, theta2, log=TRUE))
    return(likelihood+prior1+prior2)
}

# Defining initial values to be passed on to the optimizer
initVal = c(1,1)

# Finding the optimized betavector
optimResult = optim(initVal, logPosteriorMult, x=data, lambda=1, method=c(
"L-BFGS-B"),
                    control=list(fnscale=-1), lower=c(0.01,0.01), upper=c(
Inf, Inf), hessian=TRUE)

# Defining the values of interest
postMode = optimResult$par
postCov = -solve(optimResult$hessian)
names(postMode)=c("Theta1", "Theta2")
rownames(postCov)=c("Theta1", "Theta2")
colnames(postCov)=c("Theta1", "Theta2")
print("The posterior mode is:")
print(postMode)
print("The approximated standard deviation is:")
print(postCov)

## c) Discuss how a Bayesian can determine if the symmetric model in 1a) o
r the non-symmetric model in 1b)
## is most appropriate for this data. No need to compute anything here, ju
st discuss.

## By calculating marginal Likelihood for each model and check which has t
he highest. One can also calculate
## the bayes factor or the posterior model probabilities and choose the mo
del with the highest probability.

```

Assignment 2 – Bayes Linear Regression, HPD interval, predictive distribution house prices for specific house type

```

## a) Use conjugate priors, standard normal and invchisq and use BayesLinR
eg to simulate 5000 draws from posterior
## distrib

```

```

# Reading the data from file
library(MASS)
BostonHousing = Boston
y = BostonHousing$medv
X = cbind(1,BostonHousing[,1:13]) # Adding a column of ones for the interc
ept
names(X)[1] <- "intercept"
covNames <- names(X)
y <- as.numeric(y)
X <- as.matrix(X)

library(mvtnorm)

```

```

# Defining a function that simulates from the scaled inverse Chi-square di
# tribution
rScaledInvChi2 <- function(n, df, scale){
  return((df*scale)/rchisq(n,df=df))
}

BayesLinReg <- function(y, X, mu_0, Omega_0, v_0, sigma2_0, nIter){
  # Direct sampling from a Gaussian linear regression with conjugate prior
  :
  #
  #  $\beta \mid \sigma^2 \sim N(\mu_0, \sigma^2 \text{inv}(\Omega_0))$ 
  #  $\sigma^2 \sim \text{Inv-}\chi^2(v_0, \sigma^2_0)$ 
  #
  # Author: Mattias Villani, IDA, Linköping University. http://mattiasvillani.com
  #
  # INPUTS:
  # y - n-by-1 vector with response data observations
  # X - n-by-nCovs matrix with covariates, first column should be ones i
  f you want an intercept.
  # mu_0 - prior mean for beta
  # Omega_0 - prior precision matrix for beta
  # v_0 - degrees of freedom in the prior for sigma2
  # sigma2_0 - location ("best guess") in the prior for sigma2
  # nIter - Number of samples from the posterior (iterations)
  #
  # OUTPUTS:
  # results$betaSample - Posterior sample of beta. nIter-by-nCov
  s matrix
  # results$sigma2Sample - Posterior sample of sigma2. nIter-by-1 ve
  ctor

  # Compute posterior hyperparameters
  n = length(y) # Number of observations
  nCovs = dim(X)[2] # Number of covariates
  XX = t(X)%*%X
  betaHat <- solve(XX, t(X)%*%y)
  Omega_n = XX + Omega_0
  mu_n = solve(Omega_n, XX%*%betaHat + Omega_0%*%mu_0)
  v_n = v_0 + n
  sigma2_n = as.numeric((v_0*sigma2_0 + ( t(y)%*%y + t(mu_0)%*%Omega_0%*%m
  u_0 - t(mu_n)%*%Omega_n%*%mu_n))/v_n)
  invOmega_n = solve(Omega_n)

  # The actual sampling
  sigma2Sample = rep(NA, nIter)
  betaSample = matrix(NA, nIter, nCovs)
  for (i in 1:nIter){

    # Simulate from  $p(\sigma^2 \mid y, X)$ 
    sigma2 = rScaledInvChi2(n=1, df = v_n, scale = sigma2_n)
    sigma2Sample[i] = sigma2

    # Simulate from  $p(\beta \mid \sigma^2, y, X)$ 

```

```

    beta_ = rmvnorm(n=1, mean = mu_n, sigma = sigma2*invOmega_n)
    betaSample[i,] = beta_

  }
  return(results = list(sigma2Sample = sigma2Sample, betaSample=betaSample
))
}

mu_0=rep(0,ncol(X))
omega_0=1/100*diag(ncol(X))
v_0=1
sigma2_0=36
nIter=5000
post_distrib=BayesLinReg(y,X, mu_0, omega_0, v_0, sigma2_0, nIter)
post_beta=post_distrib$betaSample
colnames(post_beta)=covNames
lstat_post=subset(post_beta, select="lstat")
par(mfrow=c(1,1))
plot(density(lstat_post), main="Posterior density of lstat", lwd=2)
credInterval=quantile(lstat_post, probs=c(0.05, 0.95))
abline(v=credInterval[1], col="grey", lwd=3, lty=3)
abline(v=credInterval[2], col="grey", lwd=3, lty=3)

# Since posterior of beta is the student t-distrib the distrib is symmetric
# and therefore HPD interval is the same
# as equal tail interval

new_obs=X[9,]
names(new_obs)=covNames
new_obs_2=new_obs
new_obs_2[which(names(new_obs)=="lstat")]=new_obs_2[which(names(new_obs)=="lstat")]*0.7
post_sigma2=post_distrib$sigma2Sample
pred_price1=post_beta%%new_obs+rmnorm(nIter, mean=0, sd=sqrt(post_sigma2))
pred_price2=post_beta%%new_obs_2+rmnorm(nIter, mean=0, sd=sqrt(post_sigma2))
hist(pred_price1, breaks=50, main="Histogram of predicted price before change")
hist(pred_price2, breaks=50, main="Histogram of predicted price after change")
pred_price_house9=post_beta[,14]*(new_obs[14]*0.7-new_obs[14])
mean(pred_price_house9)
quantile(pred_price_house9, probs=c(0.025, 0.975))

# For a house like number 9 it will increase the house price with high posterior probability.

```

Assignment 4 – Simulation of predDraw from normal, approximated probability weight larger than value, decision making (linear loss function)

```

## a) Simulate 1000 draws from predictive distrib of the maximal weight on
a given future day, model:  $y=10*a$  where
##  $y$  is the weight and  $a$  is the build cost.  $y \sim N(\theta, \sigma^2)$ . Noninformative prior assumed.

```

```

y=c(191, 196, 197, 189)
sigma2=10^2
# Noninformative prior assumed to be constant

yPred_post=rnorm(1000, mean=mean(y), sd=sqrt(sigma2*(1+1/length(y))))

## b) Use simulation to approximate the predictive probability that weight
higher than 230

pred_max365=rep(0,1000)
for (i in 1:1000) {
  pred_max365[i]=max(rnorm(365, mean=mean(y), sd=sqrt(sigma2*(1+1/length(y)
))))))
}
prob_yPred365=sum(pred_max365>230)/1000
print(prob_yPred365)

## The probability is 0.157

## c) The loss function is linear

expectedLoss = function(a, maxWeight) {
  probCollapse=sum(maxWeight>10*a)/1000
  return(a*(1-probCollapse)+probCollapse*(a+100))
}

a=seq(20,30,0.01)
plot(a, sapply(a, expectedLoss, maxWeight=pred_max365), type="l", lwd=2, x
lab="a", ylab="EL",
      main="Loss function")
aOpt=a[which(min(sapply(a, expectedLoss, maxWeight=pred_max365)))]
print(aOpt)

## The answer is 23.89

```

2018-06-01

Assignment 1 – Plot posterior using samples and expression, simulate from predictive

## a) Draw 1000 samples from prior (Gamma) and 1000 samples from posterior (Gamma). Plot prior and posterior using  
## both samples and their analytical expressions.

```

n=50
x_mean=10
beta=2
nDraws=1000

```

# We know that posterior distribution is the  $\text{Gamma}(\alpha + \sum(\text{data}), \beta + n)$ . Mean for Gamma distrib is  $\alpha/\beta$ .  
## If  $\beta=2$  then  $\beta+n$  for posterior is 52.  $\alpha/2=(\alpha+500)/52$  which yields  $50*\alpha=1000$  and  $\alpha=20$   
## Check:  $20/2=10$ ,  $(20+500)/52=10$  OK!



```

alpha=20 # According to motivation above
post_draws=rgamma(nDraws, alpha+n*x_mean, beta+n)
prior_draws=rgamma(nDraws, alpha, beta)
gridWidth=0.01
muGrid_post=seq(7,12, gridWidth) # Range taken with inspiration from histogram
muGrid_prior=seq(4,20,gridWidth)
par(mfrow=c(2,1))
hist(post_draws, breaks=50, main="Posterior", xlab=expression(mu),
      freq=FALSE)
lines(muGrid_post, dgamma(muGrid_post, alpha+n*x_mean, beta+n), lwd=2, xlab=expression(mu))
hist(prior_draws, breaks=50, main="Prior", xlab=expression(mu),
      freq=FALSE)
lines(muGrid_prior, dgamma(muGrid_prior, alpha, beta), lwd=2, xlab=expression(mu))

## As seen in the plots the distributions resemble each other.

## b) Simulate 1000 draws from predictive distribution of new observation and plot distribution.

par(mfrow=c(1,1))
x_pred=rpois(1000, lambda=post_draws)
hist(x_pred, breaks=50, main="Histogram, approximated posterior predictive distribution", xlab=expression(mu),
      freq=FALSE)

## c) Prob that x51=10 based on posterior predictive distribution

sum(x_pred==10)/nDraws

## [1] 0.132

```

Assignment 2 – Bayes Linear Regression (fish), marginal posterior, 90 % equal tail with interpretation, new experiment with two different new\_obs with likelihoods (Bayesian analysis)

```

## Linear regression model for fish with 3 covariates.

# Reading the data from file
load(file = 'fish.RData')

library(mvtnorm)

# Defining a function that simulates from the scaled inverse Chi-square distribution
rScaledInvChi2 <- function(n, df, scale){
  return((df*scale)/rchisq(n,df=df))
}

BayesLinReg <- function(y, X, mu_0, Omega_0, v_0, sigma2_0, nIter){
  # Direct sampling from a Gaussian linear regression with conjugate prior
  :
  #

```

```

# beta | sigma2 ~ N(mu_0, sigma2*inv(Omega_0))
# sigma2 ~ Inv-Chi2(v_0,sigma2_0)
#
# Author: Mattias Villani, IDA, Linkoping University. http://mattiasvillani.com
#
# INPUTS:
# y - n-by-1 vector with response data observations
# X - n-by-nCovs matrix with covariates, first column should be ones if you want an intercept.
# mu_0 - prior mean for beta
# Omega_0 - prior precision matrix for beta
# v_0 - degrees of freedom in the prior for sigma2
# sigma2_0 - location ("best guess") in the prior for sigma2
# nIter - Number of samples from the posterior (iterations)
#
# OUTPUTS:
# results$betaSample - Posterior sample of beta. nIter-by-nCovs matrix
# results$sigma2Sample - Posterior sample of sigma2. nIter-by-1 vector

# Compute posterior hyperparameters
n = length(y) # Number of observations
nCovs = dim(X)[2] # Number of covariates
XX = t(X)%*%X
betaHat <- solve(XX,t(X)%*%y)
Omega_n = XX + Omega_0
mu_n = solve(Omega_n,XX%*%betaHat+Omega_0%*%mu_0)
v_n = v_0 + n
sigma2_n = as.numeric((v_0*sigma2_0 + ( t(y)%*%y + t(mu_0)%*%Omega_0%*%mu_0 - t(mu_n)%*%Omega_n%*%mu_n))/v_n)
invOmega_n = solve(Omega_n)

# The actual sampling
sigma2Sample = rep(NA, nIter)
betaSample = matrix(NA, nIter, nCovs)
for (i in 1:nIter){

  # Simulate from p(sigma2 | y, X)
  sigma2 = rScaledInvChi2(n=1, df = v_n, scale = sigma2_n)
  sigma2Sample[i] = sigma2

  # Simulate from p(beta | sigma2, y, X)
  beta_ = rmvnorm(n=1, mean = mu_n, sigma = sigma2*invOmega_n)
  betaSample[i,] = beta_

}
return(results = list(sigma2Sample = sigma2Sample, betaSample=betaSample))
}

## a) Plot marginal posterior for each param

```

```

y=as.matrix(subset(fish, select="length"))
X=as.matrix(fish[,2:ncol(fish)])
covNames=colnames(X)
mu_0=rep(0,3)
omega_0=0.01*diag(1,3)
v_0=1
sigma2_0=100^2
nIter=5000

linPost=BayesLinReg(y, X, mu_0, omega_0, v_0, sigma2_0, nIter)
betaPost=linPost$betaSample
colnames(betaPost)=covNames
sigma2Post=linPost$sigma2Sample
par(mfrow=c(2,2))
for (i in 1:ncol(betaPost)) {
  hist(betaPost[,i], xlab=paste("Beta",i,sep=""), main=paste("Marginal posterior distribution of beta", i, sep=""))
}
hist(sigma2Post, xlab=expression(sigma), main="Marginal posterior distribution of sigma2")

par(mfrow=c(1,1))

## Construct 90 % equal tail interval for beta1 and interpret it.

quantile(subset(betaPost, select="age"), probs=c(0.05, 0.95))

## It can be concluded that when the age of the fish increases with one unit the length of the fish increases
## with approximately between 2.284 and 2.960 mm with 90 % posterior probability.

## d) New experiment fish has been grown in water tank with water temp 30 degrees celsius. Newborn fish have
## have been inserted into the tank at two time points, 30 days ago and 100 days ago. Equal amount of fish
## in the two different ages. You pick up fish randomly from water tank. Do bayesian analysis (using sim methods)
## to determine predictive distrib of the length of the picked up fish.

x1=c(1,30,30)
x2=c(1,100,30)
x_pred=rep(0,nIter)
for (i in 1:nIter) {
  prob=runiform(1)
  if(prob>0.5) {
    x_pred[i]=betaPost[i,]*x1+rnorm(1, mean=0, sd=sqrt(sigmaPost[i]))
  } else {
    x_pred[i]=betaPost[i,]*x2+rnorm(1, mean=0, sd=sqrt(sigmaPost[i]))
  }
}
hist(x_pred, main="Histogram of predictive distribution of length of fish",
      xlab="Length in mm", freq=FALSE, breaks=50)

```

### Assignment 3 – Choosing between 3 models, marginal likelihood

*## c) Choose between three models where two of them use Beta prior and the last one assumes  $p=0.5$ . Which model should be chosen?*

```
model1=choose(10,3)*gamma(4)*gamma(8)*gamma(2)/gamma(12)
model2=choose(10,3)*gamma(7)*gamma(11)*gamma(8)/(gamma(4)*gamma(4)*gamma(18))
model3=choose(10,3)*0.5^10
model1_norm=model1/sum(c(model1, model2, model3))
model2_norm=model2/sum(c(model1, model2, model3))
model3_norm=model3/sum(c(model1, model2, model3))
```

### Assignment 4 – Truncated normal distrib, stan with time series model, plot of data & posterior mean & 95 % credible intervals over time

*## a) Consider observations with values above 200. Remaining datapoints asumed to be indep. and follow a truncated normal distribution with density specified. L=200 lower truncation point. Write a function in R that computes the (unnormalized) log posterior distribution of mu. Use function to plot the posterior distrib of mu for the observations greater than 200 in the data vector sulfur. For the plot, use a grid constructed in R with seq(100,400,1)*

```
# Reading the data from file
load(file = 'sulfur.RData')
```

```
muGrid=seq(100,400,1)
sigma=100
data=sulfur[sulfur>200]
```

*# Constant prior for mu is assumed*

```
logPost = function(data, mu, sigma, L=200) {
  nominator=dnorm((data-mu)/sigma, mean=0, sd=1, log=TRUE)
  denominator=log(sigma)+log(1-pnorm((L-mu)/sigma))
  return(sum(nominator-denominator+0)) # Assumed constant prior which can be set to 1 which in log scale is 0
}
```

```
post_mu=exp(sapply(muGrid, logPost, data=data, sigma=sigma))
post_mu_norm=post_mu/sum(post_mu) # Since gridwidth is 1 we don't have to compensate for it
plot(muGrid, post_mu_norm, type="l", lwd=2, main="Posterior distribution of mu", xlab=expression(mu))
```

```
library(rstan)
T = length(sulfur)
T_cens = sum(sulfur <= 200)
censData <- list(T=T, T_cens = T_cens, x=sulfur, L=200)
```

*# Model*

```

censModel <- '
data {
  int<lower=0> T;          // Total number of time points
  int<lower=0> T_cens;     // Number of censored time points
  real x[T];              // Partly censored data
  real<upper=max(x)> L;    // Lower truncation point
}

parameters {
  real mu;
  real<lower=0> sigma;
  real<upper=L> x_cens[T_cens]; // Censored values
}

model {
  int t_cens = 0;
  for (t in 1:T){
    if (x[t] > L)
      x[t] ~ normal(mu,sigma);
    else {
      t_cens += 1;
      x_cens[t_cens] ~ normal(mu,sigma);
    }
  }
}
'

```

*## b) Now consider all data points. Values below 200 being censored.*

```

fit=stan(model_code=censModel, data=censData)
print(fit)
post_draws=extract(fit)
grid=seq(1,4000,1)
plot(grid, post_draws$mu, type="l", main="Traceplot of mu", xlab=expression(mu), ylab="Value")
plot(grid, post_draws$sigma, type="l", main="Traceplot of sigma", xlab=expression(sigma), ylab="Value")
par(mfrow=c(4,2))
for (i in 1:8) {
  plot(grid, post_draws$x_cens[,i], type="l", main=paste("Traceplot of ", i, "th obs of obs below 200", sep=""),
       xlab=i, ylab="Value")
}
par(mfrow=c(1,1))

plot(post_draws$mu, post_draws$sigma, type="p", col="grey", main="Joint posterior of mu and sigma",
     xlab=expression(mu), ylab=expression(sigma))

```

*## c) Instead consider time series model. Assume that observations follow an independent normal distrib*  
*## when conditioned on a latent AR(1) process z, but with values of xi below 200 being censored and set to 200.*  
*## Modify the stan code in order to do inference for this model instead. A*

*Also put a normal prior on  
 ##  $\mu \sim N(300, 100^2)$  Plot the posterior of  $\phi$ . Also produce a plot that contains both the data and the posterior  
 ## mean and 95 % credible intervals for the latent intensity  $z$  over time.*

```
StanModel_AR = '
data {
  int<lower=0> T;          // Total number of time points
  int<lower=0> T_cens;     // Number of censored time points
  real x[T];              // Partly censored data
  real<upper=max(x)> L;    // Lower truncation point
}

parameters {
  real mu;
  real phi;
  real<lower=0> sigma;
  real<upper=L> x_cens[T_cens]; // Censored values
  vector[T] z;
}

model {
  // Prior
  int t_cens = 0;
  phi ~ uniform(-1,1);
  mu ~ normal(300, 100);
  for (n in 2:T)
    z[n] ~ normal(mu + phi * (z[n-1]-mu), sigma);

  // Model/likelihood
  for (t in 1:T){
    if (x[t] > L)
      x[t] ~ normal(z[t],20);
    else {
      t_cens += 1;
      x_cens[t_cens] ~ normal(z[t],20);
    }
  }
}

generated quantities {
  vector[T] post_mean;
  post_mean = z;
}
'

fitAR=stan(model_code=StanModel_AR, data=censData)
print(fitAR)
post_draws_AR=extract(fitAR)
postPhi=post_draws_AR$phi
postZ=post_draws_AR$post_mean
hist(postPhi, breaks=50, main="Approximated posterior density of phi", xlab=expression(phi), freq=FALSE)
grid=seq(1,31)
```

```

plot(grid, sulfur, col="blue", main="Emissions of sulfur dioxide", xlab="Day of month", ylab="mg/Nm^3",
      ylim=c(0,500))
postMean=rep(0,ncol(postZ))
credIntervals=matrix(0,ncol(postZ),2)
for (i in 1:ncol(postZ)) {
  postMean[i]=mean(postZ[,i])
  credIntervals[i,]=quantile(postZ[,i], probs=c(0.025, 0.975))
}
lines(grid, postMean, type="l", col="red", lwd=2)
lines(grid, credIntervals[,1], col="grey", lwd=1, lty=2)
lines(grid, credIntervals[,2], col="grey", lwd=1, lty=2)
legend("topleft", legend=c("Data", "Posterior mean", "95 % cred intervals"),
      lwd=c(NaN, 2, 1), lty=c(NaN,1,2),
      pch=c(1,NaN, NaN), col=c("blue", "red", "grey"))

```

2019-08-21

Assignment 1 – Bayes Linear Regression, point estimates, 95 % tail intervals, posterior mode and HPD 90 % intervals, predictive Bayesian analysis

```

## a) Use BayesLinReg to sim 5000 draws from posterior distrib of all coef
f coefficients. Summarize posterior
## with point estimate under quadratic loss function and 95 % equal tail i
ntervals. Interpret cred intervals for
## regression coefficient on nitrogen oxides concentration.

```

```

#####
##### Problem 1 #####
#####

```

*# Reading the data from file*

```

library(MASS)
BostonHousing = Boston
y = BostonHousing$medv
X = cbind(1,BostonHousing[,1:13]) # Adding a column of ones for the interc
ept
names(X)[1] <- "intercept"
covNames <- names(X)
y <- as.numeric(y)
X <- as.matrix(X)
XNewHouse <- c(1,0.03,40,1.5,0,0.5,6,30,5,3,300,17,390,4)

```

```

if(length((grep("mvtnorm",installed.packages()[,1]))))==0)
  install.packages("mvtnorm")
library(mvtnorm)

```

*# Defining a function that simulates from the scaled inverse Chi-square di*  
*stribution*

```

rScaledInvChi2 <- function(n, df, scale){
  return((df*scale)/rchisq(n,df=df))
}

```

```

BayesLinReg <- function(y, X, mu_0, Omega_0, v_0, sigma2_0, nIter){
  # Direct sampling from a Gaussian linear regression with conjugate prior

```



```

:
#
# beta | sigma2 ~ N(mu_0, sigma2*inv(Omega_0))
# sigma2 ~ Inv-Chi2(v_0,sigma2_0)
#
# Author: Mattias Villani, IDA, Linkoping University. http://mattiasvillani.com
#
# INPUTS:
# y - n-by-1 vector with response data observations
# X - n-by-nCovs matrix with covariates, first column should be ones if you want an intercept.
# mu_0 - prior mean for beta
# Omega_0 - prior precision matrix for beta
# v_0 - degrees of freedom in the prior for sigma2
# sigma2_0 - location ("best guess") in the prior for sigma2
# nIter - Number of samples from the posterior (iterations)
#
# OUTPUTS:
# results$betaSample - Posterior sample of beta. nIter-by-nCovs matrix
# results$sigma2Sample - Posterior sample of sigma2. nIter-by-1 vector

# Compute posterior hyperparameters
n = length(y) # Number of observations
nCovs = dim(X)[2] # Number of covariates
XX = t(X)%*%X
betaHat <- solve(XX,t(X)%*%y)
Omega_n = XX + Omega_0
mu_n = solve(Omega_n,XX%*%betaHat+Omega_0%*%mu_0)
v_n = v_0 + n
sigma2_n = as.numeric((v_0*sigma2_0 + ( t(y)%*%y + t(mu_0)%*%Omega_0%*%mu_0 - t(mu_n)%*%Omega_n%*%mu_n))/v_n)
invOmega_n = solve(Omega_n)

# The actual sampling
sigma2Sample = rep(NA, nIter)
betaSample = matrix(NA, nIter, nCovs)
for (i in 1:nIter){

  # Simulate from p(sigma2 | y, X)
  sigma2 = rScaledInvChi2(n=1, df = v_n, scale = sigma2_n)
  sigma2Sample[i] = sigma2

  # Simulate from p(beta | sigma2, y, X)
  beta_ = rmvnorm(n=1, mean = mu_n, sigma = sigma2*invOmega_n)
  betaSample[i,] = beta_

}
return(results = list(sigma2Sample = sigma2Sample, betaSample=betaSample))
}

```

```

mu_0=rep(0, ncol(X))
omega_0=1/10^2*diag(ncol(X))
v_0=1
sigma2_0=5^2
nIter=5000
linPost=BayesLinReg(y, X, mu_0, omega_0, v_0, sigma2_0, nIter)
betaPost=linPost$betaSample
sigma2Post=linPost$sigma2Sample
results=matrix(0,ncol(X)+1,3)
results_names=covNames
results_names=append(results_names, "sigma2")
rownames(results)=results_names
colnames(results)=c("Point estimator", "2,5%", "97,5%")
for (i in 1:ncol(X)) {
  results[i,1]=mean(betaPost[,i])
  results[i,-1]=quantile(betaPost[,i], probs=c(0.025, 0.975))
}
results[(ncol(X)+1),1]=mean(sigma2Post)
results[(ncol(X)+1),-1]=quantile(sigma2Post, probs=c(0.025, 0.975))
results

```

*## b) Kernel density estimates. Compute posterior mode and HPD 90 % for sigma2*

```

sigma2_kernel=density(sigma2Post)
sigma2_kernel.df=data.frame(sigma2=sigma2_kernel$x, density=sigma2_kernel$y)
sigma2_kernel.df=sigma2_kernel.df[order(-sigma2_kernel.df[,2]),]
index=dim(sigma2_kernel.df)[1]
sigma2_kernel.df$density=cumsum(sigma2_kernel.df$density)/sum(sigma2_kernel.df$density)
sigma2Cred=sigma2_kernel.df[sigma2_kernel.df$density<0.9,]
credInterval=c(min(sigma2Cred$sigma2), max(sigma2Cred$sigma2))
sigma2Mode=sigma2_kernel.df[1,]$sigma2

plot(sigma2_kernel, type="l", lwd=2, main="Kernel density estimate of sigma2",
      xlab=expression(sigma^2))
abline(v=sigma2Mode, col="red", lwd=1, lty=2)
abline(v=credInterval[1], col="grey", lwd=1, lty=3)
abline(v=credInterval[2], col="grey", lwd=1, lty=3)
legend("topright", legend=c("Kernel density estimate", "Posterior mode", "90 % HPD Interval"),
      lty=c(1,2,3),
      lwd=c(2,1,1), col=c("black", "red", "grey"))

```

*## c) Construction company planning to build a new house with covariates given in XNewHouse. Cost is 20000 dollars*  
*## and the company is planning to sell the house when finished. Do Bayesian analysis to determine how probable*  
*## it is that the company will make money (that the house will sell for more than 20000 dollars).*

```

XNewHouse <- c(1,0.03,40,1.5,0,0.5,6,30,5,3,300,17,390,4)
profitVec=rep(0,nIter)
for (i in 1:nIter) {

```

```

    profitVec[i]=-20+betaPost[i]*%*%XNewHouse+rnorm(1, mean=0, sd=sqrt(sigma2
Post[i]))
}
hist(profitVec)
probProfit=sum(profitVec>0)/nIter
print(probProfit)
quantile(profitVec, probs=c(0.025, 0.975))

## Very probable that the company will make a profit since 98.82 % of the
posterior draws are above zero. Negative
## values are also not present in the 95 % equal tail interval which also
indicates that the company will make
## a profit.

```

Assignment 2 – Predictive draw earthquakes (math calculations as basis)

*## b) Simulate predictive draw of max no. of years until next earthquake occurs, 95 % prob. alpha=1, beta=1.*

```

alpha=1
beta=1
xObs=c(35, 14, 4, 10, 2)
n=length(xObs)
nIter=5000
predDistrib=rep(0,nIter)
for(i in 1:nIter) {
  posteriorDraw=rbeta(1,alpha+n, beta=sum(xObs))
  predDistrib[i]=rgeom(1,posteriorDraw)
}
predDistrib_maxYear=quantile(predDistrib, probs=0.95)
predDistrib_maxYear

```

Assignment 3 – Calc unnormalized posterior and plot normalized posterior

*## c) Calc unnormalized posterior and plot normalized posterior. Gamma prior and indep likelihoods.*

```

gridWidth=0.01
thetaGrid=seq(0,2,gridWidth)
xData <- c(1.888, 2.954, 0.364, 0.349, 1.090, 7.237)
yData <- c(-1.246, -1.139, -0.358, -1.308, -0.930, -0.157, -0.111, -0.635)
alpha=3
beta=2

logPosteriorX = function(theta, alpha, beta) {
  return(dgamma(theta, alpha, beta, log=TRUE))
}

likeY = function(y, theta) {
  return(-3*sum(log(1+(1/5)*(y-log(theta))^2)))
}

logPosterior = function(theta, alpha, beta, xDat, yDat) {
  likelihoodY=likeY(yDat, theta)
  logPostX=logPosteriorX(theta, length(xDat)+3, sum(xDat)+2)
  return(likelihoodY+logPostX)
}

```

```

}

post_theta=sapply(thetaGrid, logPosterior, alpha=alpha, beta=beta, xDat=xData, yDat=yData)
post_theta_norm=1/gridWidth*exp(post_theta)/sum(exp(post_theta))
plot(thetaGrid, post_theta_norm, type="l", lwd=2, main="Posterior of theta", xlab=expression(theta), ylab="Density")

```

Assignment 4 – Simulate using RWM Hastings, simulation using metropolis hastings with gamma proposal density function, suggestions how to improve sampler, traceplots

*## Aircraft incidents assumed to be independent, follow negative binomial distrib. Assume joint prior*

*##  $1/\phi^2$*

*## a) Simulate from posterior using Metropolis algorithm. Denote  $\theta=c(m, \phi)$  and use as proposal dens*

*## the multivariate normal density (random walk metropolis).*

*# Load airline incidents data*

```
load(file = 'incidents.RData')
```

```
data=incidents$incidents
```

```
library(mvtnorm)
```

```
nIter=1000
```

```
burnIn=50
```

```
theta_0=c(200,20)
```

```
c=0.1
```

```
postCov=diag(c(100,5))
```

*# Defining function for sampling through metropolishastings*

```

RVMSampler = function(previousVal, postCov, c, myFunction, ...) {
  proposalVal=rmvnorm(1, mean=previousVal, sigma=c*postCov)
  proposalVal[proposalVal<=0]=1e-6
  alpha=min(1, exp(myFunction(proposalVal,...)-myFunction(previousVal, ...)))
  u=runif(1)
  if(u < alpha) {
    return(list(theta=proposalVal, acceptProb=alpha))
  } else {
    return(list(theta=previousVal, acceptProb=alpha))
  }
}

```

```
logPrior = function(phi) {
```

```
  return(-2*log(phi))
```

```
}
```

```
logLike <- function(param, x){
```

```
  theta1 = param[1]
```

```
  theta2 = param[2]
```

```
  logPost = sum(logdNegBin(x, theta1, theta2)) - 2*log(theta2)
```

```
  return(logPost)
```

```
}
```

```

logPost = function(theta, data) {
  log_Prior=logPrior(theta[2])
  log_Like=logLike(theta, data)
  return(log_Prior+log_Like)
}

post_matrix = matrix(0, nIter+burnIn, 2)
# Setting initial values of beta to same initVals as in the optimizer (taken randomly from normal distrib)
post_matrix[1,]=theta_0
accProb=rep(0, nIter)
set.seed(12345)

for(i in 1:(nIter+burnIn)) {
  if(i<(nIter+burnIn)) {
    draw=RVMsampler(post_matrix[i,], postCov, c, logPost, data)
    post_matrix[i+1,]=draw$theta
    accProb[i+1]=draw$acceptProb
  }
}

iter=seq(1,nIter+burnIn,1)
plot(iter[-(1:burnIn)], post_matrix[-(1:burnIn),1], type="l", lwd=1, col="grey", main="Traceplot of mu in RVM",
      xlab=expression(mu), ylab="Value")
plot(iter[-(1:burnIn)], post_matrix[-(1:burnIn),2], type="l", lwd=1, col="grey", main="Traceplot of phi in RVM",
      xlab=expression(phi), ylab="Value")
mean(accProb)

## This MCMC sampler is not efficient since it moves very slowly and is therefore probably not exploring
## the whole posterior distribution. We can also see that the acceptance probability for this algorithm
## is around 84,4 % and it should be around 30 %. Once could tune the c parameter to lower the acceptance probability.
## One example is to increase c to a value of 3 which would yield in approximately 30 % acceptance rate.

## b) Instead simulate from posterior using metropolis hasting.

c=0.8

MHSampler = function(previousVal, postCov, c, myFunction, ...) {
  proposalVal_mu=rgamma(1, c*previousVal[1], c)
  proposalVal_phi=rgamma(1, c*previousVal[2], c)
  proposalVal=c(proposalVal_mu, proposalVal_phi)
  proposalVal[proposalVal<=0]=1e-6
  alpha=min(1, exp(myFunction(proposalVal,...)-myFunction(previousVal, ...))
  +
              dgamma(previousVal[1], c*proposalVal[1], c)+dgamma(previousVal[2],c*proposalVal[2],c)-
              dgamma(proposalVal[1], c*previousVal[1], c)-dgamma(pr

```

```

oposalVal[2], c*proposalVal[2],c)))
  u=runif(1)
  if(u < alpha) {
    return(list(theta=proposalVal, acceptProb=alpha))
  } else {
    return(list(theta=previousVal, acceptProb=alpha))
  }
}

post_matrix2 = matrix(0, nIter+burnIn, 2)
theta_0=c(200,10)
post_matrix2[1,]=theta_0
accProb2=rep(0, nIter)
set.seed(12345)

for(i in 1:(nIter+burnIn)) {
  if(i<(nIter+burnIn)) {
    draw=MHSampler(post_matrix2[i,], postCov, c, logPost, data)
    post_matrix2[i+1,]=draw$theta
    accProb2[i+1]=draw$acceptProb
  }
}

plot(iter[-(1:burnIn)], post_matrix2[-(1:burnIn),1], type="l", lwd=1, col=
"grey", main="Traceplot of mu in MH",
      xlab=expression(mu), ylab="Value")
plot(iter[-(1:burnIn)], post_matrix2[-(1:burnIn),2], type="l", lwd=1, col=
"grey", main="Traceplot of phi in MH",
      xlab=expression(phi), ylab="Value")
mean(accProb2)

## The new algorithm seems to rapidly explore the posterior which is good.
The acceptance probability is also lower
## around 30 % which also indicates that this algorithm is better than the
previous one.

```

2019-10-31

Assignment 1 – Expected utility, calculations from math

*## a) Theta is known*

```

theta=0.6
eu_buy=0.6*30-0.4*10
eu_nobuy=0.6*90-120*0.4

```

*## Answer: Should buy option*

*## c) Compute bayesian decision for day 101 based on information in b).*

```

theta_new=13/21
eu_buy_new=theta_new*30-(1-theta_new)*10
eu_nobuy_new=theta_new*90-(1-theta_new)*120

```

*## Answer: Should buy since utility higher.*

Assignment 2 – Poisson likelihood with gamma prior, plot posterior, separation of data, now two indep poisson models, comparison

*## a) Consider poisson likelihood model. Use conjugate prior and plot posterior in given interval.*

*## Compute posterior probability that theta is smaller than 21.*

*# Calculations show that alpha=20, beta=1*

```
data=Traffic$y
alpha=20
beta=1
n=length(data)
```

*# We know that Poisson with gamma prior is gamma distributed with alphaNew=alpha+sum(data), betaNew=beta+n*

```
grid=seq(18,24,0.01)
post_distrib=rgamma(grid, shape=alpha+sum(data), rate=beta+n)
plot(grid, post_distrib, type="l", lwd=2, main="Posterior distrib. of theta",
      xlab=expression(theta))
post_prob=pgamma(21, shape=alpha+sum(data), rate=beta+n)
```

*## Answer: Probability is 0.0557*

*## b) Two independent poisson models.*

```
data_model1=Traffic[which(Traffic[,3]=="yes"),]$y
data_model2=Traffic[which(Traffic[,3]=="no"),]$y
```

```
alpha_1=20+sum(data_model1)
alpha_2=20+sum(data_model2)
beta_1=1+length(data_model1)
beta_2=1+length(data_model2)
post_distrib_1=rgamma(5000, shape=alpha_1, rate=beta_1)
post_distrib_2=rgamma(5000, shape=alpha_2, rate=beta_2)
hist(post_distrib_1, breaks=50)
hist(post_distrib_2, breaks=50)
post_diff=post_distrib_2-post_distrib_1
hist(post_diff,
      main="Posterior distribution of difference between no speedlimit and speedlimit",
      xlab="No. of accidents")
quantile(post_diff, prob=c(0.025, 0.975))
mean(post_diff)
```

*## We can see that the difference between the two distributions is larger than 0 with high probability. In this*

*## case we can say that the difference in traffic accidents between when no speed limit were applied and*

*## when a speed limit were applied is between 2.82 and 5.53 approximately with 95 % posterior probability.*

*## The conclusion from this is that yes, a speed limit leads to a lower amount of accidents.*



```
## c) A politician claims that the experiment proves that introducing speed limit decreases the number of accidents by at least 15 %.
```

```
mean(0.85*post_distrib_2>post_distrib_1)
```

```
## Likely that the decrease yields 15 % but 86 % probable and not 95 % probability which is commonly used  
## in statistical experiments.
```

Assignment 3 – Simulation of joint posteriors with Gibbs sampling, traceplots

```
## c) Make simulations of joint posterior of  $\nu$  and  $\pi$  using Gibbs sampling.
```

```
x=20  
lambda=10  
alpha=2  
beta=2  
nIter=2000  
burnIn=500  
  
results=matrix(0,burnIn+nIter,2)  
initVal=lambda # Since lambda=30  
results[1,1]=initVal  
results[1,2]=rnorm(1)  
for (i in 1:(nIter+burnIn-1)) {  
  z=rpois(1, lambda*(1-results[i,2]))  
  results[i+1,1]=z+x  
  results[i+1,2]=rbeta(1, alpha+x, beta+results[i+1,1]-x)  
}  
  
grid=seq(burnIn+1, nIter+burnIn)  
barplot(table(results[(burnIn+1):(nIter+burnIn),1]), main="Marginal posterior of  $\nu$ ", xlab=expression(nu))  
hist(results[(burnIn+1):(nIter+burnIn),2], breaks=50, main="Marginal posterior of  $\pi$ ", xlab=expression(pi))  
plot(grid, results[(burnIn+1):(nIter+burnIn),2], type="l")  
plot(grid, results[(burnIn+1):(nIter+burnIn),1], type="l")  
  
## Convergence seems good since markov chain is exploring full posterior and have good mixing.
```

Assignment 4 – Stan, plot scatter plot & mean of posterior predictive distrib & 90 % equal tail intervals, stanmodel with heteroscedastic variance (different variance over time)

```
## a) Use supplied stan model to do Bayesian inference. Draw 2000 posterior samples and use 500 for burnin.  
## Produce figure with scatter plot, overlay curve for mean of posterior predictive distrib, in range [0,25].  
## Also overlay curves 90 % equal tail interval for same posterior predictive distrib given values of x in range [0,25]
```

```
# Load data  
cars = cars
```

```

library(rstan)
LinRegModel <- '
data {
  int<lower=0> N;
  vector[N] x;
  vector[N] y;
}
parameters {
  real alpha;
  real beta;
  real<lower=0> sigma2;
}
model {
  sigma2 ~ scaled_inv_chi_square(5,10);
  for (n in 1:N)
    y[n] ~ normal(alpha + beta * x[n], sqrt(sigma2));
}
'

x=cars$speed
y=cars$dist
nIter=2000
burnIn=500
N=dim(cars)[1]
data=list(N=N,x=x,y=y)
fit=stan(model_code=LinRegModel, data=data, iter=nIter, warmup = 500, chains=1)
print(fit)
postDraws=extract(fit)
alpha_draws=postDraws$alpha
beta_draws=postDraws$beta
sigma_draws=postDraws$sigma2
xGrid=seq(0,25)
n=length(alpha_draws)
mean_credInt=matrix(0,length(xGrid),3)
count=1
for (i in 1:length(xGrid)) {
  ysim=rep(0,length(nIter-burnIn))
  ysim=alpha_draws+beta_draws*xGrid[i]+rnorm(nIter-burnIn, mean=0, sd=sqrt(sigma_draws))
  mean_credInt[count,1]=mean(ysim)
  mean_credInt[count,-1]=quantile(ysim, probs = c(0.05, 0.95))
  count=count+1
}

plot(x,y,xlab="Speed", ylab="Distance", col="blue", main="Plot for model with constant sigma prior")
lines(xGrid, mean_credInt[,1], lwd=2, col="red")
lines(xGrid, mean_credInt[,2], lwd=1, lty=2)
lines(xGrid, mean_credInt[,3], lwd=1, lty=2)
legend("topleft", legend=c("Data", "Posterior mean", "90 % cred interval"), col=c("blue", "red", "grey"), pch=c(1, NaN, NaN), lty=c(NaN, 1, 2), lwd=c(NaN, 2, 1))

```

*## b) Compute 95 % equal tail credible interval for alpha. Give real-world interpret of the interval.*

```
quantile(alpha_draws, probs=c(0.025, 0.975))
```

*## The interpretation of the credible interval for alpha is that if the car has no speed it travels a negative distance between -31 and 4.25 approximately with 95 % posterior probability. This is not realistic. To prevent this a prior can be set to alpha with a mean around zero which however would make the linear prediction worse. One can also use the Log Normal distribution for y to force it to have a value above zero.*

*## c) Reproduce results in b) with heteroscedastic variance.*

```
LinRegModel_hetero <- '
data {
  int<lower=0> N;
  vector[N] x;
  vector[N] y;
}
parameters {
  real alpha;
  real beta;
  real<lower=0> sigma2[N];
  real gamma;
  real phi;
}
model {
  for (n in 1:N)
    sigma2[n] ~ scaled_inv_chi_square(5,exp(gamma+phi*x[n]));
  for (n in 1:N)
    y[n] ~ normal(alpha + beta * x[n], sqrt(sigma2[n]));
}
'

data=list(N=N,x=x,y=y)
fit2=stan(model_code=LinRegModel_hetero, data=data, iter=nIter, warmup = 500, chains=1)
print(fit2)
postDraws2=extract(fit2)
alpha_draws=postDraws2$alpha
beta_draws=postDraws2$beta
sigma_draws=postDraws2$sigma2
xGrid=seq(0,25)
n=length(alpha_draws)
mean_credInt=matrix(0,length(xGrid),3)
count=1
for (i in 1:length(xGrid)) {
  rinvs=rchisq(nIter-burnIn, 5)
  sigma_draw=5*exp(postDraws2$gamma + xgrid[i] * postDraws2$phi)^2/rinvs
  ysim=rep(0,length(nIter-burnIn))
  ysim=alpha_draws+beta_draws*xGrid[i]+rnorm(nIter-burnin, mean=0, sd=sqrt(sigma_draw))
}
```

```

mean_credInt[count,1]=mean(ysim)
mean_credInt[count,-1]=quantile(ysim, probs = c(0.05, 0.95))
count=count+1
}

plot(x,y,xlab="Speed", ylab="Distance", col="blue", main="Plot of model wi
th heteroscedastic sigma prior")
lines(xGrid, mean_credInt[,1], lwd=2, col="red")
lines(xGrid, mean_credInt[,2], lwd=1, lty=2)
lines(xGrid, mean_credInt[,3], lwd=1, lty=2)
legend("topleft", legend=c("Data", "Posterior mean", "90 % cred interval")
, col=c("blue", "red", "grey"),
      pch=c(1, NaN, NaN), lty=c(NaN, 1, 2), lwd=c(NaN, 2, 1))

## The new model seems to capture the data better than the old one.

```